

# Reuse, but Verify: Certified Maintenance of LLM-Computed Table Cells under Prompt Edits

Arjun Lohan  
Independent Researcher  
United States  
arjunlohan7@gmail.com  
lohan@usc.edu

## Abstract

AI-native tables let users define columns in natural language: each cell is computed by a language model from a versioned prompt over its row. When the prompt is edited, deployed systems face a choice: recompute every cell, which is expensive and rate-limited, or silently reuse stale ones, which incurs unbounded error. We present sIVM, a maintenance layer that certifies which cached cells may be reused after a prompt edit, bounding the false-reuse rate of each certified stratum by a user budget  $\alpha$  with probability at least  $1 - \delta$ . sIVM freezes cached-value strata, samples adaptively on a doubling look schedule, and bounds each stratum’s flip rate with an exact finite-population confidence bound; a strict mode certifies at a deflated threshold to bound the error of the reused cells alone. We also establish an impossibility floor for this estimand: a model’s rate of disagreement with itself under the unedited prompt floors the certifiable budget. On two public corpora (89,184 Stack Overflow developer profiles; 210,250 Djinni candidate CVs) with 28,000 labeled cells and a production table system, a formatting-only edit certifies 88.4% call savings at  $\alpha=0.2$  with 3.20% realized false-reuse, and at deployment scale 90 oracle calls certify reuse of 81,424 of 89,184 cells. Across 396,000 calibration runs with planted flip rates, unsafe certificates occur in at most 1.00% of runs in any configuration, against a nominal per-stratum budget of 5%. A second certification 55 days later, against the same cache, measured provider-side model drift as a higher flip rate and re-certified the same budgets, the tighter one a look deeper than on the August draws. Certification reproduces on 6 of 6 further model families, while which edits are benign varies by family. We release the system under a source-available noncommercial license, with the labels and a benchmark of versioned prompt-edit pairs.

## 1 Introduction

Interactive artificial intelligence (AI) tables turn a natural-language prompt into a *materialized view*: a column  $V = (\pi, M, \theta)$  with prompt  $\pi$ , model  $M$ , and pinned (held fixed) decode parameters  $\theta$ , materialized as cells  $c_i = f_V(r_i)$  over rows  $r_1, \dots, r_n$ . The category is now a product surface rather than a research proposal: Google Sheets evaluates an  $\text{AI}()$  formula per row and offers AI columns in tables [16], Airtable computes AI fields per record [1], Notion autofills database properties with a model [33], and Snowflake Cortex applies  $\text{AI\_COMPLETE}$  and  $\text{AI\_CLASSIFY}$  to table columns in SQL [43] (Microsoft’s Excel COPILOT function shipped as a preview and was withdrawn in September 2026, a

measure of how quickly this surface is moving [32]). These views are edited constantly: users iterate on prompts the way they iterate on spreadsheet formulas. Every edit invalidates the cache by construction, and today’s systems invalidate it wholesale: an edit to an 89,184-row column triggers 89,184 fresh model calls, or the user silently keeps stale values with no bound on how wrong they are.

*Running example.* One column from our evaluation recurs throughout. Over 89,184 Stack Overflow (SO) developer profiles, a Boolean AI column asks “*Is this person a data-platform specialist?*”, judged from each profile’s role, databases, platforms, and tools. Its owner edits the prompt to insert two paragraph breaks and lowercase one word, changing no content word (the *formatting-only* edit of §5; Table 10 shows every edit we study). Recomputing costs 89,184 model calls; keeping every cached value silently carries the edit’s flip rate, 6.3% on this pair, into the table. sIVM instead groups the cached cells by their cached value (the two *strata*, cached-TRUE and cached-FALSE), recomputes a small seeded sample from each, and reuses a stratum only when a statistical bound on its flip rate clears the user’s error budget: on this column at deployment scale, 90 fresh calls certified reuse of 81,424 of 89,184 cells, and the same 90-call budget did so again 55 days later against a model that had drifted in the meantime.

Recent semantic data systems bound the error of *one-shot* query execution against a gold model (cascades, semantic operators, approximate processing with guarantees), and verified semantic caches bound reuse under a *fixed* definition. Neither answers the maintenance question an interactive table poses: *after this specific edit  $\pi \rightarrow \pi'$ , which cached cells are still valid, and with what guarantee?*

### 1.1 Why intuition fails here

Three measurements reframe the problem (§5). First, large language model (LLM) cells are stochastic even at temperature zero [4, 35]: the same prompt on the same row disagrees with itself 5.0% of the time on our Boolean workload (within-version draw pairs on  $n=2,000$ ; an earlier  $n=200$  two-draw probe read 3.0%) and 9.5% at default temperature, so “did the edit change this cell?” is only meaningful relative to a noise floor, the rate at which the model disagrees with itself under the unedited prompt. Second, edit direction is not monotone: under a strictly scope-*widening* edit, 45.7% of cached-TRUE cells flipped ([39.0, 52.7] over 199 such cells), refuting the folk lemma that widening preserves positives. Third, an edit’s surface class is not a safe predictor of its semantic impact. On the same column, a meaning-preserving synonym paraphrase flipped 8.90% of cells where a formatting-only edit

flipped 6.3%. On a candidate-profile column from a second corpus, a formatting-only edit flipped 12.30% against 12.45% for a deliberate criteria change, rates that are indistinguishable. Neither pair orders the way the edits’ labels suggest. These are five author-written edits, one per class per column, so they show that class can fail to predict flip rate, not how often it does; the certifier below never consults the class, and §5 shows the same edit’s flip rate varying widely across model families. Both comparisons hold the column fixed deliberately: across columns, the column’s own stability dominates the edit.

Sound reuse therefore has to be *measured*: sample a few fresh cells, statistically bound the flip rate of the rest, and reuse only what clears the user’s error budget.

## 1.2 Contributions

- (1) **Problem formalization** (§2): certified cross-version reuse for stochastic prompt-defined views, with a pair of estimands (the population quantities a guarantee is about) that are well-defined under decode pinning: the realized false-reuse rate over the cells a certified stratum presents, which we certify by default, and over the reused cells alone, which the strict mode certifies at a deflated threshold; and its impossibility floor (the intrinsic flip rate).
- (2) **The sIVM procedure** (*semantic incremental view maintenance*; §3): frozen cached-value stratification (with optional interaction refinement), grid-adaptive sampling on a doubling schedule of *looks* (points at which the sample is inspected; deciding at several looks is the *peeking* the validity theorem must survive, handled by splitting  $\delta$  across looks), futility stopping, per-stratum certification with an exact finite-population confidence bound, and two guarantee targets (presented cells; strict reuse set). Any frozen stratifier affects power only, never validity, where power is the chance of certifying a stratum that does meet the budget. The bound is a pluggable component: the exact bound is proven for the without-replacement sampling the procedure performs, and four alternatives (two empirical-Bernstein bounds, the binomial Clopper–Pearson bound, and a betting confidence sequence) are measured inside the same procedure.
- (3) **A production implementation** (§4) inside a working AI-table system (per-cell runner, content-hash cache, cost ledger, certificates, the persisted records of what cleared the budget, surfaced in the user interface (UI)), released under a source-available noncommercial license, not a simulation harness.
- (4) **An empirical anatomy** (§5) on two real corpora with 28,000 labeled cells: the noise floor, the edit-class savings curve, the  $\alpha$ -cost frontier, minority-stratum volatility and its vote-of- $k$  diagnosis, a stratifier ablation showing that refinement is validity-free but costs power on this benchmark, a calibration study that measures the rate of unsafe certificates against the nominal budget, a check of the cross-row independence assumption, a model-family study on 6 further families, and a re-certification 55 days

**Table 1: Notation.**

| Symbol                       | Meaning                                                                                           |
|------------------------------|---------------------------------------------------------------------------------------------------|
| $V = (\pi, M, \theta)$       | a view: prompt template, model, pinned decode parameters; $V' = (\pi', M, \theta)$ after the edit |
| $r_i, c_i$                   | row $i$ and its cached cell, one draw from $f_V(r_i)$                                             |
| $X'_i, F_i, p_i$             | a fresh draw under $V'$ ; the flip indicator $\mathbb{1}[c_i \neq X'_i]$ ; its mean               |
| $\equiv$                     | the column’s equivalence relation (exact match here)                                              |
| $j, n_j, s_j$                | a stratum, its size, its self-flip rate (floor)                                                   |
| $\mathcal{S}, m_j$           | the sampled (freshly recomputed) cells; those in stratum $j$                                      |
| $\mathcal{R}$                | the reuse set: cached cells kept under a certificate                                              |
| $L_j, t$                     | stratum $j$ ’s look schedule and a look in it                                                     |
| $u_{j,t}$                    | the upper confidence bound on stratum $j$ ’s flip rate at look $t$                                |
| $\alpha, \delta$             | error budget on the false-reuse rate; probability that any certificate is wrong                   |
| $K, \delta_j$                | number of strata; the per-stratum budget $\delta/K$ , split evenly over looks                     |
| $\text{FR}_j^{\text{pres}}$  | realized false-reuse of stratum $j$ over the cells it presents                                    |
| $\text{FR}_j^{\text{reuse}}$ | the same over its reused cells alone                                                              |

after the deployment run that measures provider-side drift (a higher flip rate against the same cache) and shows the certifier pricing it: the budgets certified in August certify again, the tight one a look deeper than on the August draws, and the tightest budget, on which the August draws are silent, is refused.

- (5) **A reusable benchmark** (§6): versioned prompt-edit pairs with full ground-truth flip labels over public corpora redistributable under the Open Database License (ODbL) and the MIT License, enabling like-for-like evaluation of future maintenance systems.

## 2 Problem

*Terminology.* Four words recur. A *stratum* is a group of cached cells fixed before any sampling (by default, the cells that share a cached value). A *look* is a point in the sampling schedule at which the sample drawn so far is inspected and a decision may be taken. A stratum’s *floor* is its self-flip rate, the rate at which the model disagrees with its own earlier answer when nothing has changed; no budget below the floor is certifiable. A *certificate* is the persisted record that a stratum cleared the budget, with the evidence it cleared it on. A few statistical terms recur as well. An *oracle draw* (or oracle call) is one fresh model call under the edited prompt for one row. An *estimand* is the population quantity a guarantee is about. *Power* is the chance that a stratum which does meet the budget is in fact certified. *Pinned* means held fixed for the duration of a certification (the model, the decode parameters, the stratification). The *Bonferroni* split divides an error budget among several tests so that their failures cannot add up past it. Table 1 collects the symbols.

*Views and cells.* An AI column is a view  $V = (\pi, M, \theta)$ : a prompt template  $\pi$  with typed output contract, a model  $M$ , and pinned decode parameters  $\theta$ . For row  $r_i$ , the bound prompt  $\pi(r_i)$  interpolates only the fields  $\beta(\pi)$  that  $\pi$  references; the stored cell  $c_i$  is *one draw* from the output distribution  $f_V(r_i)$ . Even at temperature zero this distribution is not degenerate: providers are nondeterministic across calls, as prior measurements of hosted models document [4, 35], in part because the floating-point reduction order inside serving kernels changes with the co-batched load [18, 47], so every quantity below is defined over draw randomness.

*The maintenance event.* The user edits the definition,  $V \rightarrow V' = (\pi', M, \theta)$ . For each cached cell define the flip indicator  $F_i = \mathbb{1}[c_i \neq X'_i]$ , where  $X'_i \sim f_{V'}(r_i)$  is a fresh draw under the new definition and  $\equiv$  is the column’s equivalence relation (exact match for the Boolean and select columns studied here).  $F_i$  is Bernoulli with unknown mean  $p_i = \Pr[c_i \neq X'_i]$ .  $p_i$  deliberately folds together the edit’s semantic effect and decode noise on both sides: it answers the operational question of whether recomputing cell  $i$  now would return a value different from the one shown.

*Estimand and guarantee.* A maintenance policy partitions the cached cells into a sampled set  $\mathcal{S}$  (freshly recomputed), a reuse set  $\mathcal{R}$  (stale values kept), and a recompute set. Two error rates matter, and they differ. The *presented-cells* rate of a certified stratum  $j$  is

$$\text{FR}_j^{\text{pres}} = \frac{1}{n_j} \sum_{i \in \mathcal{S}} F_i :$$

of all  $n_j$  cells the user now sees from that stratum, the fraction that would flip if recomputed (sampled cells are fresh and contribute zero); a flip is any disagreement with one fresh draw, decode noise included, as the previous paragraph defines it. That zero is a convention of the single-draw estimand, not a property of recomputed cells: a recomputed cell is itself one draw, and a further draw disagrees with it at the stratum’s self-flip rate (24.1% and 3.1% on our Boolean column, §5). The estimand charges reuse for exactly what the user would see change on recomputation, which is the operational question, and the impossibility floor below is the price of that choice: an excess-over-floor target,  $p_j - s_j$ , would certify the identity edit for free, and we deliberately do not certify it, because a user who recomputes sees one draw and lives with it. The *reuse-set* rate

$$\text{FR}_j^{\text{reuse}} = \frac{1}{n_j - m_j} \sum_{i \in \mathcal{S}} F_i$$

normalizes by the  $n_j - m_j$  reused cells alone and is never smaller. A certificate bounds the stratum’s whole flip count, and hence  $\text{FR}_j^{\text{pres}}$ , but not  $\text{FR}_j^{\text{reuse}}$ : the remainder is adversely selected, because certification happens exactly when the sample looked clean (on the headline pair at  $\alpha=0.2$  the two read 3.20% and the reuse-set rate whose interval is [2.5, 4.2], Table 2). sIVM therefore certifies the presented-cells estimand by default, and offers a strict reuse-set mode that certifies at the deflated threshold  $\alpha(n_j - m_j)/n_j$ , which bounds  $\text{FR}_j^{\text{reuse}}$  by conservation of flips. Theorem 3.1 states both guarantees for the realized rates of this run, not for their expectations. Both realized rates are reported for every experiment (Table 2). All guarantees are per-stratum and therefore hold simultaneously for every union of certified strata, not only the aggregate (baseline B2 in §5, aggregate-only certification, shows why the aggregate alone is operationally misleading).

*Reading the budget.* A certificate says: with probability at least  $1 - \delta$  over this run’s sampling and oracle draws, at most an  $\alpha$  fraction of the certified stratum’s cells would flip if recomputed (default mode), or at most an  $\alpha$  fraction of its reused cells alone (strict mode).  $\alpha$  is that share;  $\delta$  is the probability that any certificate of the run is wrong at all. Both are the user’s to set, and the right values are workload-specific: a recruiter screen that a person reviews row by row tolerates a looser  $\alpha$  than a column that feeds a downstream aggregate, and the strict mode exists for the latter. The realized error under a certificate is typically far below  $\alpha$  (3.20% against  $\alpha=0.2$  on the headline pair, Table 2), because the bound must cover the worst population compatible with the sample; the frontier of §5 shows what tightening  $\alpha$  costs in certifiable savings, and the floor below shows where tightening it becomes impossible.

#### Assumptions.

*Assumption 1 (Pinned snapshot).*  $M$  and  $\theta$  are fixed during certification; oracle draws are distributed as  $f_{V'}(r_i)$  throughout the certification window and are independent *across rows* (provider-side batching that correlates concurrent responses [18, 47] would violate this; our runner issues independent requests, and §5, “Independence across rows,” compares sequential with concurrent requests). The validity proof uses independence only through non-interference: a sampled cell’s draw must not depend on which other cells were sampled or in flight. Certificates are scoped to the snapshot the oracle draws came from; the cache may be older. A provider-side model change ends the snapshot, and §5 measures one (change-of-world maintenance is complementary work).

*Assumption 2 (Frozen stratification).* The stratifier is a deterministic function of the cache, the edit, and row content, fixed before any oracle call. Any such stratifier preserves validity; its quality affects only power (§5, “The stratification hierarchy”).

*Assumption 3 (Uniform within-stratum sampling).* Sampled rows are drawn by a seeded uniform shuffle within each stratum, without replacement, and each sampled row receives one fresh oracle draw, independent across rows (Assumption 1). The calibration study of §5 probes this assumption where it matters, at planted rates just above the budget.

*No monotonicity shortcuts.* Folk reasoning suggests one-sided safety: a scope-widening edit should preserve TRUE cells, so cached TRUEs could be reused for free. Our measurements refute the unconditional form: under a widening edit, 45.7% of cached-TRUE cells flipped, driven largely by boundary instability, not semantics. Edit direction therefore enters sIVM only as a stratification hint; any directional certificate must be *conditional* on a measured violation rate, never assumed.

*The impossibility floor is per-stratum.* Let  $s_j$  be stratum  $j$ ’s self-flip rate: the probability that two independent draws of the *same* cell disagree, *conditional on the cached draw that put the cell in stratum  $j$* . For the identity edit  $\pi'=\pi$ ,  $p_j$  equals  $s_j$ , so no procedure whose oracle is a single fresh draw can certify stratum  $j$  at  $\alpha < s_j$ : the certifiable budget is floored by model stability. Two scoping notes apply. The conditioning is not incidental. Because membership is decided by one draw, a minority-value stratum is selected partly *for* having flipped that way, so  $s_j$  is a property of the (cache,

stratifier) pair, not of the column, and recomputing the cache with vote-of- $k$  at write time lowers it (§5, “Minority-stratum anatomy”). And the edit-independence is derived for the identity edit only; an edit that sharpens the output distribution can drive  $p_j$  below  $s_j$ , so we state edit-independence as an empirical regularity on our workloads, not as a bound.

The floors are unequal, and the column-level average describes neither stratum. On our Boolean column the cached-FALSE stratum has  $s=3.1\%$  (95% row-cluster bootstrap [2.5, 3.7]) while the cached-TRUE stratum has  $s=24.1\%$  ([19.6, 28.8]), against a pooled 5.0% that is their mixture. Rows are bucketed by the stored cached value the certifier itself stratifies on, and the interval resamples rows, since draw-pairs from one row are dependent. Three otherwise separate observations follow from this single measurement. The FALSE stratum is the only one that certifies at all, and it is certifiable in principle down to budgets near its own floor; at  $\alpha=0.05$  it certifies in 74.4% of replications, at the deepest look of the schedule, for 38.6% mean savings (Table 2): there, power rather than the floor is what binds. The TRUE stratum is never certified at any budget we tried: its floor exceeds every budget we ran. Its size sets the deployment-scale savings ceiling: with 7,715 of 89,184 cells cached TRUE, reuse cannot exceed 91.3%. The ceiling belongs to *this* cache, and a cache built with vote-of- $k$  at write time would hand the certifier a smaller stratum. And any method that reuses cells from that stratum inherits an error rate at or above its floor no matter what aggregate target it was given, which is what we observe when we run one (§5, B2’). Column-level stability is the wrong unit; the certifier’s stratification recovers the right one, and a standing refusal on a stratum is a stability diagnosis of that slice of the column.

### 3 The sIVM procedure

Given the edit  $\pi \rightarrow \pi'$  and cached cells for the affected rows (Figure 1):

- (1) **Diff.** Compute a word-level longest-common-subsequence (LCS) diff of the prompts; collect content terms of added/removed spans (the *delta*). The delta is recorded in the certificate for provenance and consumed only by the optional refinement in step 2; it never enters the reuse decision itself.
- (2) **Stratify (frozen).** Partition rows by cached value. We adopt the skeleton of proxy-certified selection [22] deliberately here: with a binary proxy this step is their reduction, transported from one-shot selection to maintenance. What we change is the guarantee it carries, from one aggregate bound over the selected set to a bound per stratum, which is the only difference §5 (B2’) isolates and the only one that protects a subgroup whose floor differs from the column’s. Optionally refine each value stratum by an interaction bucket (lexical delta-term overlap or embedding cosine between the delta text and the row’s bound content, tertile-bucketed): a validity-free refinement that pays only when within-value flip risk is heterogeneous, and otherwise costs power through the  $\delta/K$  split (§5).
- (3) **Schedule.** For stratum  $j$  of size  $n_j$ , fix a doubling look schedule  $L_j = (45, 90, 180, \dots)$  capped at  $n_j$ , with at most six looks; split the stratum’s error budget  $\delta_j = \delta/K$  evenly

---

#### Algorithm 1 sIVM certification of an edit $\pi \rightarrow \pi'$

---

**Require:** cached cells  $\{(r_i, c_i)\}$ ; edit  $\pi \rightarrow \pi'$ ; budget  $\alpha$ ; confidence  $\delta$ ; first look  $n_0$ ; at most  $\ell$  looks; mode  $\in \{\text{presented, reuse-set}\}$   
**Ensure:** reuse set  $\mathcal{R}$ , recompute set, certificate record

- 1:  $\Delta \leftarrow$  word-level LCS diff of  $\pi$  and  $\pi'$
- 2: strata  $1, \dots, K \leftarrow$  partition of the cells by cached value, optionally refined by an interaction bucket of  $\Delta$ ; freeze
- 3: **for** each stratum  $j$  with  $n_j$  cells **do**
- 4:    $\text{order}_j \leftarrow$  seeded shuffle of the stratum’s rows in canonical order;  $L_j \leftarrow (n_0, 2n_0, \dots)$  capped at  $n_j$ , at most  $\ell$  looks;  $\delta' \leftarrow \delta/(K |L_j|)$
- 5:   **for** each look  $n \in L_j$  **do**
- 6:     draw  $X'_i$  for the first  $n$  rows of  $\text{order}_j$  not yet drawn;  $k \leftarrow \{i \leq n : c_i \neq X'_i\}$
- 7:      $u \leftarrow$  exact upper bound of Eq. (1) at  $(k, n, n_j, \delta')$
- 8:      $\tau \leftarrow \alpha$  if presented mode, else  $\alpha(n_j - n)/n_j$
- 9:     **if**  $n < n_j$  and  $u \leq \tau$  **then**
- 10:       certify stratum  $j$ ; **break**
- 11:     **end if**
- 12:     **if**  $k/n > \alpha$  **then**
- 13:       refuse stratum  $j$  (futility); **break**
- 14:     **end if**
- 15:   **end for**
- 16:   refuse stratum  $j$  if its schedule is exhausted
- 17: **end for**
- 18:  $\mathcal{R} \leftarrow$  unsampled cells of certified strata; recompute set  $\leftarrow$  unsampled cells of refused strata
- 19: persist the certificate (budget,  $\Delta$ , per-stratum evidence) and apply  $\mathcal{R}$  additively

---


across looks. Every experiment here pins  $n_0=45$ , at most six looks,  $\delta=0.1$ , and  $K=2$  value strata, so a look on a large stratum is tested at  $\delta' = 0.1/12$ ; a small stratum has fewer looks and a correspondingly larger per-look level.

- (4) **Sample and test.** Order the stratum’s rows canonically, shuffle them once with the seed, and at each look  $t$  reveal fresh draws for the prefix. Compute the exact finite-population upper confidence bound  $u_{j,t}$  of Eq. (1) on the stratum flip rate at level  $\delta_j/|L_j|$ ; *certify and stop* if  $u_{j,t} \leq \alpha$ ; *stop for futility* if the empirical rate alone exceeds  $\alpha$ .
- (5) **Apply.** Reuse all unsampled rows of certified strata (recording certificate provenance per cell); recompute the rest. Sampled rows are already fresh.

Algorithm 1 states the procedure in full.

*The bound.* Fix a stratum of  $N$  cells ( $N = n_j$ ), and give every cell one hypothetical fresh draw under  $V'$ , the coupling on which the presented-cells estimand is defined; let  $M$  be the number of cells whose fresh draw flips. The sample of  $n$  cells is drawn uniformly without replacement and independently of those draws, and shows  $k$  flips. Conditional on the  $N$  outcomes,  $k$  is Hypergeometric( $N, M, n$ ), and  $\Pr[k' \leq k \mid M]$  is non-increasing in  $M$ . The exact upper confidence bound at level  $\delta'$  is

$$u = \frac{1}{N} \max\{M : \Pr[\text{Hyp}(N, M, n) \leq k] > \delta'\}, \quad (1)$$



**Figure 1: The procedure on the configuration our experiments pin.** Stratification is frozen before any oracle call, so it moves power but never validity (Assumption 2); only steps 3–4 spend model calls, and they stop on evidence about a stratum’s rate, which is independent of the stratum’s size. The two lanes are the outcome we measure throughout: the stable stratum certifies and is reused under a certificate, while the volatile one is refused at every budget we tried and recomputed. A refusal is a stability diagnosis of that slice of the column (§2), reached at the cost of one look instead of a full recomputation.

the finite-population form of the Clopper–Pearson bound [12], to which it tends as  $N$  grows. Because the coverage holds conditionally on every population realization, it holds unconditionally over sampling and draw randomness together, and it bounds the realized flip count  $M$ , not only its expectation. This is the natural bound for a binary loss under the sampling the procedure performs. On a clean sample its cost is  $\ln(1/\delta')/n$  to first order; the empirical-Bernstein bounds pay range terms with constants of  $7/3$  (Maurer–Pontil, on  $\ln(2/\delta')$ ) and  $7/3 + 3/\sqrt{2}$  (Bardenet–Maillard, on  $\ln(5/\delta')$ ) that a binary loss does not need. On a clean sample of  $n=90$  from a stratum of 1,800 cells at the pinned per-look level, the exact bound reads 0.050 (0.052 in the 81,469-cell deployment stratum, where it coincides with the binomial limit to three digits), the Maurer–Pontil bound [31] that an earlier version of this article certified with reads 0.144, and the Bardenet–Maillard without-replacement bound [6] reads 0.317 (§5 measures the difference on every configuration and prices a betting confidence sequence [46] and time-uniform alternatives [20]). The system exposes the alternatives as ablation arms.

**THEOREM 3.1 (VALIDITY UNDER PEEKING).** *Let  $K$  be the number of strata produced by the frozen stratifier,  $\delta_j = \delta/K$ , and let every look use the exact bound of Eq. (1) at level  $\delta_j/|L_j|$ . Under Assumptions 1–3, with probability at least  $1 - \delta$  over sampling and oracle draws, every certified stratum  $j$  has realized whole-stratum flip count  $M_j \leq \alpha n_j$ ; consequently  $\text{FR}_j^{\text{pres}} \leq \alpha$  for every certified stratum and for every union of certified strata (the empty certification is vacuously valid). In reuse-set mode, where certification at sample size  $m_j$  requires  $u_{j,t} \leq \alpha(n_j - m_j)/n_j$ , the same event implies  $\text{FR}_j^{\text{reuse}} \leq \alpha$ . The statements hold for any per-look bound that is a valid upper confidence bound at that level for uniform without-replacement draws, including the Bardenet–Maillard bound [6]; the Maurer–Pontil bound [31] is stated for independent draws and is retained only as an ablation arm.*

**PROOF.** Fix stratum  $j$  and condition on its  $n_j$  fresh-draw outcomes, of which  $M_j$  flip. If  $M_j > \alpha n_j$ , certifying at look  $t$  requires  $u_{j,t} \leq \alpha < M_j/n_j$ , i.e., the observed count to fall in the lower tail that Eq. (1) rejects at level  $\delta_j/|L_j|$ ; by the monotonicity of the hypergeometric tail in  $M$ , that event has probability at most  $\delta_j/|L_j|$  over the sampling, for every population realization, hence unconditionally. A union bound over the  $|L_j|$  looks and the  $K$  strata gives  $\delta$ ; this is the Bonferroni instance of Learn then Test [2], each (stratum, look) pair testing the null  $M_j > \alpha n_j$  with a certificate as the rejection. The futility rule only stops sampling, never certifies, so it cannot increase the error probability. On the good event,  $\text{FR}_j^{\text{pres}} = (M_j - k_j)/n_j \leq \alpha$ , where  $k_j$  is the flips observed in the sample, and any size-weighted union of certified strata inherits the bound. In reuse-set mode the good event is  $M_j \leq \alpha(n_j - m_j)$ ; at worst every flip falls in the remainder, so  $\text{FR}_j^{\text{reuse}} \leq \alpha$ . The conditioning uses Assumption 1 only through non-interference: the population of fresh-draw outcomes must be the same whichever cells the shuffle selects.  $\square$

**Remark.** The Bardenet–Maillard bound is proven for the same uniform without-replacement draws (its Serfling factor shrinks the variance term as the sampled fraction grows; its constants are those of their Theorem 4.3), so the argument transfers to it unchanged. The Maurer–Pontil bound is stated for independent draws, and whether it remains valid at the sampled fractions our runs reach (74%) is not something we prove, which is why it is no longer the pinned bound.

**Power.** In a large stratum a clean sample certifies once  $(1-\alpha)^n \leq \delta'$ , i.e., at  $n^* = \lceil \ln \delta' / \ln(1-\alpha) \rceil$ ;  $n^*=94$  at  $\alpha=0.05$  under the pinned procedure ( $K=2$ , six looks), against 257 for the Maurer–Pontil bound. The doubling schedule cannot stop on  $n^*$ : the first look at or above it is the one it takes, so a clean stratum of  $n_j=5000$  pays 180 draws in calibration, while one of  $n_j=600$  pays 90, because its finite-population correction and its shorter schedule’s larger per-look level bring the earlier look under  $\alpha$ . Either way, savings

require strata much larger than  $n^*$ , i.e., large tables, which is precisely the deployment regime. Near the frontier, certifying a stratum with true rate  $p$  requires roughly  $n \gtrsim z_{\delta}^2 p(1-p)/(\alpha-p)^2$  samples (the normal approximation to the exact test), which explains the measured collapse of savings as  $\alpha \downarrow p$  (Table 2, SO formatting: 88.1%, 82.6%, and 38.6% mean savings at  $\alpha = 0.2, 0.1$ , and  $0.05$ , and refusal at  $0.02$ ).

## 4 Implementation

sIVM runs inside *lore*, a working AI-table system: a natural-language search bar compiles to a versioned filter specification executed identically on MySQL and Elasticsearch (conformance-tested for count, page, sort, aggregation, and hydration parity on 89,184 rows); AI columns are per-cell jobs with typed structured outputs, bounded concurrency, retry, and a content-hash cache keyed on *only the bound fields*  $\beta(\pi)$ , so unrelated row changes never invalidate the cache and content-duplicate rows share one computation. Every run passes a pre-run cost gate and lands in a ledger; prompt edits create immutable versions (the sIVM substrate); certificates persist with full strata evidence, and reused cells carry visible provenance in the UI. Decode is pinned ( $T=0$ ) for all cell computation per Assumption 1.

*Storage model.* Cells are keyed by (column, row, prompt version), so versions of a column *coexist* instead of replacing one another, and a secondary index on (column, version, model, content hash) serves the cache. This is the one design decision sIVM depends on. Systems that key a cell by row alone, or that hash the prompt into the cache key, destroy the old materialization at the moment of the edit, and the question this paper asks cannot then be posed: there is nothing left to certify. Keeping versions side by side turns “which cached cells still transfer?” into a join between two versions of one table, which is why the certifier is a query over the store, not a service beside it. Two details of that query matter for reproducibility and for Assumption 1. The certifier orders each stratum’s rows canonically (by content hash, then row identifier) before the seeded shuffle, so the sample path is a function of the seed and the data alone; an earlier version took the order the storage engine happened to return, which changed when the query planner switched indexes, and the same seed then drew a different sample (§5). And a model snapshot is a version of its own: oracle draws taken after a provider-side change are stored under a new prompt version carrying the same template, so draws from two snapshots never share a key and cannot be mixed in one certification path.

*Certificates are records, not decisions.* A certificate persists the budget  $\alpha$ , the prompt delta, the per-stratum evidence (size, cells sampled, bound attained, certified or refused), and the sampled, reused, and recompute counts. It is therefore auditable after the fact: the realized-error audits in §5 are recomputed from the certificate’s own persisted reuse set by a script that shares no code with the certifier. Applying a certificate is a single `INSERT ... SELECT` that carries certified values forward under the new version and stamps the certificate identifier on each one.

That insert carries an anti-join against the target version, so reuse is strictly *additive*: a certificate can supply a value where

none was computed, and can never overwrite one that was. An earlier implementation used an upsert, which is the obvious way to write it and is silently destructive: applying a certificate after a partial recompute reverted freshly computed cells to the stale values the certificate had been asked to reuse. The bug is invisible in aggregate error, because a reverted cell and a reused cell are the same row with the same value; it surfaces only if something outside the certifier knows what that cell should have said. Any system that maintains stochastic cells across versions has this hazard, and the guard belongs in the write path, not in code review.

## 5 Experiments

*Workloads.* All measurements use two public corpora and three prompt-defined columns (Table 10). On the SO Developer Survey (89,184 respondent rows) a Boolean column asks whether the respondent is a data-platform specialist, judged from structured facets of the row (role, databases, platforms, tools); it exists twice, as the production column and as a lab replica seeded from the production column’s second version, so that edits can be studied without touching the deployed table. On Djinni English candidate profiles (210,250 rows) a four-way ordinal select column assigns a seniority tier from the position, years of experience, and free-text profile. Five edit pairs across the three columns cover formatting-only, synonym rewording, scope widening, and criteria change; each is labeled on the same seeded  $n=2,000$  evaluation vector.

*Setup and evaluation protocol. Scope.* One budget-class reasoning model carries the full experimental program; 6 further families from other providers replicate the lab column’s two pairs, 3 of them all five pairs, and the quantitative frontier reported below is specific to the (edit, model) pair while the qualitative structure is what the family study tests (§8). Table 3 lists every model identifier and run date. *Models and decoding.* Every cell in the primary program uses DeepSeek V4 Flash (gateway identifier `deepseek/deepseek-v4-flash-0731`), called through the Vercel AI Gateway between 4 August 2026 and 5 August 2026, and again on 29 September 2026 for the re-certification and the assumption checks; the model-family study calls further families through the same gateway on 15 August 2026 to 29 September 2026. Every call passes temperature 0 and no other decode parameter (maximum output length and reasoning effort stay at the provider’s defaults; no decode seed is sent, since the seeds in this paper govern row sampling: mulberry32, main seed 42, sampling replications 1000 to 1999). Cells are requested as typed structured outputs against a JSON schema through the same client library; a call whose output fails the schema or whose request errors is retried by the client, and a cell that still fails is stored with an error status and excluded from labels and from every rate in this paper (no such cell exists among the cells written through the column runner; the probe scripts that call the model directly report their own exclusions where they occur). The embedding proxy in baselines B1 and B3 and in the interaction stratifier is `openai/text-embedding-3-small`, also through the gateway. Every ground-truth pass, sample, and replay is seeded and reproducible from released scripts. Flip labels come from single fresh draws under the new version, which matches the estimand (a label answers “would a recompute differ from the cache now?”). A three-draw relabeling

**Table 2: The pinned primary procedure (cached-value strata; exact bound; adaptive doubling looks  $n_0=45$ , six looks;  $\delta=0.1$ ; presented-cells estimand) on all five edit pairs,  $n=2,000$  each at temperature 0. Columns: *flip*, the pair’s label flip rate; *oracle* and *reused*, the calls spent and the cells reused by the main seed;  $FR_{pres}$  and  $FR_{reuse}$ , the main seed’s realized false-reuse over presented cells and over the reuse set alone (the latter with a 95% Wilson confidence interval, CI); *savings*, the mean over  $B=1,000$  sampling replications on frozen ground-truth labels with a 95% percentile interval (PI); *cert.*, the share of those replications that issued any certificate; *exceed.*, the share of the certifying replications (not of  $B$ ) whose realized reuse-set rate exceeded  $\alpha$ , a metric the default mode does not certify (§2). The table is generated mechanically from the result files; earlier exploratory sweeps with other stratifiers and schedules are reported as ablations, not headline numbers.**

| Edit                   | $\alpha$ | flip  | oracle | reused | $FR_{pres}$ | $FR_{reuse}$ [95% CI] | savings % [95% PI] | cert. | exceed. |
|------------------------|----------|-------|--------|--------|-------------|-----------------------|--------------------|-------|---------|
| SO formatting-only     | 0.02     | 6.3%  | 225    | 0      | –           | –                     | 0.0 [0.0,0.0]      | 0%    | –       |
| SO formatting-only     | 0.05     | 6.3%  | 1485   | 374    | 0.55%       | 2.67 [1.5,4.9]        | 38.6 [0.0,81.7]    | 74%   | 0.0%    |
| SO formatting-only     | 0.10     | 6.3%  | 90     | 1769   | 3.20%       | 3.28 [2.5,4.2]        | 82.6 [72.7,88.4]   | 99%   | 0.0%    |
| SO formatting-only     | 0.20     | 6.3%  | 90     | 1769   | 3.20%       | 3.28 [2.5,4.2]        | 88.1 [86.2,88.4]   | 100%  | 0.0%    |
| SO synonym rewording   | 0.02     | 8.9%  | 90     | 0      | –           | –                     | 0.0 [0.0,0.0]      | 0%    | –       |
| SO synonym rewording   | 0.05     | 8.9%  | 90     | 0      | –           | –                     | 0.0 [0.0,0.0]      | 0%    | –       |
| SO synonym rewording   | 0.10     | 8.9%  | 90     | 0      | –           | –                     | 55.6 [0.0,89.0]    | 81%   | 0.0%    |
| SO synonym rewording   | 0.20     | 8.9%  | 225    | 1645   | 5.81%       | 6.44 [5.4,7.7]        | 87.4 [82.3,89.0]   | 100%  | 0.0%    |
| SO scope widening      | 0.02     | 9.4%  | 90     | 0      | –           | –                     | 0.0 [0.0,0.0]      | 0%    | –       |
| SO scope widening      | 0.05     | 9.4%  | 765    | 0      | –           | –                     | 1.0 [0.0,0.0]      | 1%    | 100.0%  |
| SO scope widening      | 0.10     | 9.4%  | 405    | 1441   | 4.50%       | 5.62 [4.5,6.9]        | 66.3 [0.0,87.8]    | 89%   | 0.0%    |
| SO scope widening      | 0.20     | 9.4%  | 90     | 1756   | 5.33%       | 5.47 [4.5,6.6]        | 86.6 [85.5,87.8]   | 100%  | 0.0%    |
| Djinni formatting-only | 0.02     | 12.3% | 180    | 0      | –           | –                     | 0.0 [0.0,0.0]      | 0%    | –       |
| Djinni formatting-only | 0.05     | 12.3% | 315    | 0      | –           | –                     | 0.0 [0.0,0.0]      | 0%    | –       |
| Djinni formatting-only | 0.10     | 12.3% | 495    | 383    | 4.17%       | 8.09 [5.8,11.3]       | 8.4 [0.0,32.6]     | 74%   | 8.4%    |
| Djinni formatting-only | 0.20     | 12.3% | 953    | 1047   | 8.33%       | 10.12 [8.4,12.1]      | 52.1 [32.6,73.6]   | 100%  | 0.0%    |
| Djinni criteria change | 0.02     | 12.4% | 180    | 0      | –           | –                     | 0.0 [0.0,0.0]      | 0%    | –       |
| Djinni criteria change | 0.05     | 12.4% | 315    | 0      | –           | –                     | 0.0 [0.0,0.0]      | 0%    | –       |
| Djinni criteria change | 0.10     | 12.4% | 540    | 378    | 3.52%       | 6.88 [4.7,9.9]        | 14.4 [0.0,32.4]    | 83%   | 5.4%    |
| Djinni criteria change | 0.20     | 12.4% | 745    | 1255   | 7.50%       | 11.00 [9.4,12.8]      | 48.6 [32.4,69.5]   | 100%  | 0.0%    |

of the headline pair quantifies the instrument’s own variance: with fresh labels on both sides the edit flips 7.2% of cells, versus 6.0% when only the v2 side is drawn fresh against the frozen evaluation vector, and 6.30% under majority-vote labels; 4.8% of rows relabel between instruments. That 6.0% is the same population quantity Table 2 reports as 6.3% for this edit, measured on the same 2,000 rows with a different v2 draw; the two differ by 7 rows, which is what single-draw labeling costs, and we report the churn instead of picking whichever instrument flatters the result. (The within-version floor, 5.0%, is the separate quantity that bounds certifiability.) Table 4 reconciles every flip-rate figure quoted for the headline pair and its cached-TRUE stratum with the edit, the cells, and the instrument behind it. Labels are individually noisy with about a point of draw-to-draw spread in the aggregate, which the replication standard deviations in Table 2 dominate, and it is population functionals, not labels, that the guarantee addresses; this is the same picture the LLM-as-annotator literature draws, where model labels reach crowd-worker or expert agreement on many tasks but with task-specific noise that has to be validated rather than assumed [15, 36, 44, 49]. Ledgered model spend across every experiment in this paper is \$12.01 for 100,683 billed cells, read from the system’s own cost ledger. That figure is a lower bound, not a total: the ledger records only cells written through the column runner, and a further 99,957 cells were computed by probe scripts that call the model directly (the self-flip probes, the vote-of-3 study, the multi-draw relabeling, the model-family runs, and the independence check).

*Baselines.* Four comparisons isolate what sIVM adds. (B1) *Verified-cache similarity has no cross-version signal:* the per-row

similarity between the v1-bound and v2-bound prompts, the quantity a verified semantic cache thresholds, predicts cell flips at chance as measured by the area under the receiver-operating-characteristic curve (AUC 0.454 formatting, 0.499 widening; on the formatting edit the score’s entire range is 0.96–0.99, because shared row text dominates). Adding “a version key” to such a cache therefore cannot produce a sound cross-version policy. Our interaction score carries real signal exactly when the edit is semantic (AUC 0.669 on widening) and none above chance when flips are noise-driven (AUC 0.406 on formatting), which is why the cached-value dimension of stratification carries the formatting case. (B2) *Aggregate-only certification (no strata)*, run with the same schedule and budgets as the pinned procedure, is valid for its aggregate estimand and silently corrupts subgroups. On the formatting edit it certifies at  $\alpha=0.1$  (82.0% savings, 6.46% aggregate error) while the reused cached-TRUE subgroup carries 35.3% error, and at  $\alpha=0.2$  it certifies 97.8% at 6.39% aggregate with 37.2% subgroup error; on the widening edit it refuses at  $\alpha=0.1$  and certifies 95.5% at  $\alpha=0.2$  with 44.6% subgroup error. Per-stratum bounds are what make the guarantee mean something for every user-visible slice of the table, and the effect is not an artifact of giving the baseline a shorter schedule. (B2’) *A SUPG-style competitor, matched.* The closest existing method is proxy-certified selection [22]: score rows with a cheap proxy, spend an oracle budget labeling a sample, threshold the proxy so the selected set’s error clears the target. We transported it to maintenance with the proxy the reduction in §3 names, the cached value itself, and matched it to sIVM on every axis we could: identical oracle budget, identical labels, identical error target, the same exact bound on the selected population, the same total error budget  $\delta$ , each split

**Table 3: Every model the article uses, by gateway identifier.**

| Gateway identifier              | Role                                                                                                 | Run dates                                         | \$/M tokens in / out |
|---------------------------------|------------------------------------------------------------------------------------------------------|---------------------------------------------------|----------------------|
| deepseek/deepseek-v4-flash-0731 | primary model: every cell of the main program, the deployment run, the drift and independence checks | 4 August 2026 to 5 August 2026; 29 September 2026 | 0.08 / 0.15          |
| openai/text-embedding-3-small   | embedding proxy (baselines B1, B3; interaction stratifier)                                           | with the main program                             | 0.02 / 0.00          |
| google/gemini-2.5-flash-lite    | family study, all five pairs                                                                         | 15 August 2026; 29 September 2026                 | 0.10 / 0.40          |
| google/gemini-3-flash           | family study, lab column's two pairs                                                                 | 15 August 2026                                    | 0.50 / 3.00          |
| zai/glm-4.7-flash               | family study, all five pairs                                                                         | 15 August 2026; 29 September 2026                 | 0.07 / 0.40          |
| alibaba/qwen3.7-flash           | family study, all five pairs                                                                         | 15 August 2026; 29 September 2026                 | 0.03 / 0.13          |
| openai/gpt-5-nano               | family study, lab column's two pairs                                                                 | 29 September 2026                                 | 0.05 / 0.40          |
| google/gemini-3.8-flash         | family study, lab pairs and part of the rest                                                         | 29 September 2026                                 | 0.75 / 3.75          |

**Table 4: Reconciliation of the flip-rate figures quoted in this article. Each is a different quantity or a different instrument on the same quantity; none is a restatement of another.**

| Figure                                                       | Edit                              | Cells                                       | Instrument (artifact)                                                                                                         |
|--------------------------------------------------------------|-----------------------------------|---------------------------------------------|-------------------------------------------------------------------------------------------------------------------------------|
| <i>Headline pair: SO formatting-only, Boolean lab column</i> |                                   |                                             |                                                                                                                               |
| 6.3%                                                         | formatting v1→v2                  | $n=2,000$ evaluation vector                 | stored v2 oracle cells against the stored v1 cache, one draw each side; the Table 2 label instrument (exp8)                   |
| 6.0%                                                         | formatting v1→v2                  | same rows                                   | a second, independent v2 draw against the same frozen v1 cache (exp10); differs from the row above by 7 rows                  |
| 6.30%                                                        | formatting v1→v2                  | same rows                                   | majority of three draws on each side (exp10)                                                                                  |
| 7.2%                                                         | formatting v1→v2                  | same rows                                   | fresh single draws on both sides, neither the cache (exp10)                                                                   |
| 5.0%                                                         | identity (no edit)                | same rows                                   | within-version draw pairs: the self-flip floor, pooled over both cached values (exp10)                                        |
| 3.1%, 24.1%                                                  | identity                          | cached-FALSE, cached-TRUE rows              | the same floor per cached-value stratum (exp10, stored v1 value)                                                              |
| <i>Cached-TRUE stratum figures</i>                           |                                   |                                             |                                                                                                                               |
| 37.1%                                                        | formatting v1→v2                  | cached-TRUE rows of the vector              | single fresh draw against the cache (exp4)                                                                                    |
| 24.7%                                                        | formatting v1→v2                  | same                                        | vote-of-3 on both sides (exp4)                                                                                                |
| 45.7%                                                        | widening v1→v2, production column | 199 cached-TRUE rows                        | single fresh draw against the cache (Wilson [39.0, 52.7])                                                                     |
| 35.3%                                                        | formatting v1→v2                  | cached-TRUE cells B2 reused at $\alpha=0.1$ | realized error among the cells the aggregate-only baseline reused, an adversely selected subset, not a population rate (exp7) |
| 37.2%                                                        | formatting v1→v2                  | same at $\alpha=0.2$                        | same (exp7)                                                                                                                   |

across the number of tests its own procedure performs (sIVM across strata and looks,  $\delta/12$  per test; the competitor across its candidate thresholds,  $\delta/3$ ), and the same reported quantity (error among reused cells, not the presented-cells rate sIVM certifies).

Matched this way it is a live competitor that wins on volume and loses on where its errors land. At  $\alpha=0.2$  on the formatting edit it reused 1,910 cells to sIVM's 1,769 from the same 90 labels, and at 6.54% realized error it satisfies the bound it was given, against sIVM's 3.28%. On the widening edit the same pattern holds: 1,910 reuses at 9.58% versus 1,756 at 5.47%. The extra volume is drawn from the volatile stratum, 183 cached-TRUE cells on the formatting edit and 196 on the widening edit, where the per-stratum floor is 24.1% and no aggregate target can make reuse safe; sIVM reuses none of them. §2 predicts this behavior: an aggregate target is satisfied by any mixture that averages out, so it buys volume with the correctness of the least stable subgroup, and per-stratum bounds are what extend the guarantee to that part of the table.

With the cached value as the proxy there are only three candidate thresholds, and at  $\alpha=0.2$  the aggregate bound clears at the loosest one, so the arm selects 100.0% of the rows it was allowed to on the formatting edit and 100.0% on the widening edit. Matched this way, the competitor reduces to B3, reuse everything, with a certificate attached: its bound is satisfied, its aggregate error is real,

and its behavior is indistinguishable from the guarantee-free status quo. The difference sIVM makes here is that per-stratum bounds refuse a subgroup an aggregate target is willing to spend.

We note two caveats on the same experiment. First, run with an embedding-interaction proxy instead of the cached value the reduction names, and with its much larger threshold search corrected accordingly, the competitor certifies nothing at any budget we tried: a proxy that does not separate flipping rows from stable ones cannot be rescued by the selection procedure wrapped around it, which is what our own stratifier ablation teaches from the other direction. Second, the competitor's per-test level is looser than sIVM's because it runs fewer tests, and its volume advantage is sensitive to that: the comparison is matched on the quantity a user sets ( $\delta$  in total), not on the level of any individual test, and we report both so the reader can see which one the result rests on.

(B3) *Guarantee-free embedding transfer* (GPTCache-style [5],  $\tau=0.95$ ) reuses every cell on the formatting edit and 84.0% of them on the widening edit, silently absorbing error at roughly the population flip rates (6.3%; 9.45%): the deployed status quo. That  $\tau$  sits on the shoulder of the widening edit's collapse, quantified in the threshold sweep below. Table 5 lines these policies up beside the two trivial ones, always recompute and always reuse, on the two headline edits at  $\alpha=0.2$ : the trivial policies bracket the



design space (zero error at full cost; full savings at the population flip rate), the aggregate methods buy volume from the unstable subgroup, and sIVM is the only row whose reused cells carry a per-stratum bound.

*Noise floors and the stability law.* Replicate self-runs give the two column types’ intrinsic self-flip rates. The Boolean column’s floor is 5.0%, measured from within-version draw pairs over  $n=2,000$  rows (an initial  $n=200$  two-draw probe read 3.0% and is superseded); the 4-way select’s floor is 6.5% ( $n=200$  probe, a smaller sample than the Boolean floor’s and stated as such; every observed self-flip was between adjacent tiers). At default temperature the Boolean floor rises to 9.5% on that same  $n=200$  probe, whose  $T=0$  arm read 3.0%: the claim rests on the within-instrument contrast, and the default-temperature floor alone exceeds a 5% budget, which motivates decode pinning. Formatting-only edit flip rates sit above each floor (6.3% over 5.0%; 12.30% over 6.5%): on the Boolean column formatting flips are floor-dominated, and on these two columns the flip rate tracks the column’s own floor more closely than the edit’s semantic class (the Djinni criteria change lands at 12.45%, indistinguishable from its formatting edit at 12.30%); five author-written pairs cannot say how general that is.

*Other model families.* Every measurement above uses one model, so we ran the lab column’s two pairs on 6 further budget-class families from different providers: same seeded rows as the primary evaluation vector, same prompts,  $T=0$ , one fresh draw per cell, the pinned procedure (Table 6). Three properties transfer across families; one does not. First, every family has a measurable self-flip floor, from 0.4% to 34.8% across them, and the floor bounds certifiability as §2 predicts: GPT-5 nano, whose floor of 34.8% sits above every budget in the grid, flips 39.6% of cells under the formatting edit, barely above that floor, and is refused on every stratum of that edit after 90 oracle calls. Second, certified reuse is not a property of the primary model. On Gemini 3 Flash (floor 0.4%, formatting flip 4.5%) the formatting edit certifies the cached-FALSE stratum at both budgets; on Qwen3.7-Flash (floor 5.5%, flip 14.3%) it certifies both strata at the loose budget after most of the schedule; and on the three families whose formatting edit is not benign, the synonym pair certifies a stratum instead: the cached-FALSE stratum on Gemini 2.5 Flash-Lite, and the cached-TRUE stratum on GLM-4.7-Flash and on GPT-5 nano, where it is the one stratum the loose budget admits. Gemini 3.8 Flash, the family added in this revision, certifies both strata at both budgets under the formatting edit, which flips 0.9% of its cells. Its floor is 0.9%. Certification reproduces on 6 of the 6 families, and in every certified cell of Table 6 the realized presented-cells error is inside the budget. Third, refusals stay conservative: a refused stratum is never reused, and where a family refuses a stratum whose flip rate is below the budget (GLM-4.7-Flash’s cached-FALSE stratum on the formatting edit, 17.0% against  $\alpha=0.2$ ), the futility rule stopped at the first look because the sample rate exceeded the budget, a power miss the frontier of Figure 2 predicts near the budget, not a validity failure. What does not transfer is which edits are benign. The formatting-only edit that is floor-dominated on the primary model (6.3% over 5.0%) flips 25.4% of cells on Gemini 2.5 Flash-Lite (floor 4.9%) and 25.0% on GLM-4.7-Flash (floor 10.3%),

in both cached-value strata (cached-TRUE and cached-FALSE flip 28.7% and 24.6% on the former, 27.0% and 17.0% on the latter). Decomposed on Gemini 2.5 Flash-Lite, the two paragraph breaks alone flip 23.2% and the one-word case change alone 11.0%, and the flips run one way (24.6% of them TRUE→FALSE on Gemini 2.5 Flash-Lite, 86.4% on GLM-4.7-Flash): a shift of the model’s decision threshold, in opposite directions on the two families, not noise. On all three families whose formatting edit fails it is the synonym rewording, semantically a tightening of the definition, that certifies a stratum, because its flips concentrate in one cached value and leave the other reusable. The label ordering that intuition suggests (formatting benign, semantic edits not) therefore inverts on some families and holds on others; benignness is a property of the (edit, model) pair, and the certifier measures it per stratum without ever consulting the label. sIVM’s value is a function of the model’s self-consistency and of the edit’s effect on that model, both of which it measures before reusing anything. The remaining three pairs (the production column’s scope widening and the two Djinni select pairs) were run on 4 of these families with the same protocol (Table 7), so that all five pairs are complete on 3 of these families (Qwen3.7-Flash, Gemini 2.5 Flash-Lite, GLM-4.7-Flash); 4 of the 4 certify at least one stratum on those pairs, and the same reading holds: which pair is benign, and on which stratum, is a property of the family, and the per-stratum floor of the select column varies across families in the same way the Boolean column’s does.

*Main results.* Table 2 reads as follows. On the stable Boolean column, the certifier reuses aggressively where safe (formatting at  $\alpha \geq 0.1$ : 88.1% mean savings, interval [86.2, 88.4], at  $\alpha=0.2$ , with realized presented-cells error 3.20% and a reuse-set rate whose Wilson interval is [2.5, 4.2], the  $FR_{\text{pres}}$  and  $FR_{\text{reuse}}$  columns of the table; 82.6% mean savings at  $\alpha=0.1$ , where the main seed certifies at its first look, 90 oracle calls, and the replications span 90–405), certifies the formatting edit at  $\alpha=0.05$  in 74.4% of replications by the deepest look, and refuses semantic edits at tight budgets; on the less stable select column it certifies at  $\alpha=0.2$ , and at  $\alpha=0.1$  in most replications, and treats formatting and criteria edits alike, as their flip rates warrant. 12 of the 20 default-mode configurations certify in at least one replication; budgets below a stratum’s intrinsic rate are refused, and the refusals are the frontier of Figure 2, not failures of the procedure. Realized presented-cells error is inside  $\alpha$  in every main-seed certification. Across the 1,000 replications of every configuration, a certificate whose realized presented-cells error exceeds  $\alpha$  occurs in 1 configuration (SO scope widening at  $\alpha=0.05$ , where the certified stratum’s flip rate sits just above the budget); it certifies in 1.2% of replications and 75% of those exceed, an unconditional rate of 0.90% against the  $\delta=0.1$  the run was given. That is an exact test spending its budget at a tight null, the behavior the calibration study below measures at scale, not a defect; the Maurer–Pontil bound of the earlier version never produced such a cell because it left almost all of  $\delta$  unspent (§5, “The bound”). The reuse-set metric evaluated on presented-mode certifications exceeds  $\alpha$  in 3 frontier cells (SO scope widening at  $\alpha=0.05$  (100.0% of the 1% of replications that certify); Djinni formatting-only at  $\alpha=0.1$  (8.4% of the 74% of replications that certify); Djinni criteria change at  $\alpha=0.1$  (5.4% of the 83% of replications that certify)). This does not violate  $\delta$ , which bounds the presented-cells event, nor Theorem 1, whose

**Table 5: Reuse policies on the SO formatting and scope-widening edits at  $\alpha=0.2$  ( $n=2,000$ , main seed). Realized error is the flip rate among reused cells; the trivial policies are exact by construction (always-reuse error is the population flip rate). B2 additionally reports the error of its reused cached-TRUE subgroup in parentheses. B2’ is matched to the 90 labels sIVM spent; B2 runs the pinned schedule without strata and stopped after 45 labels on the formatting edit (360 at  $\alpha=0.1$ ) and 90 on the widening edit.**

| Policy                    | Guarantee   | Formatting |               | Widening |                  |
|---------------------------|-------------|------------|---------------|----------|------------------|
|                           |             | reused     | error         | reused   | error            |
| Always recompute          | exact       | 0.0%       | 0.0%          | 0.0%     | 0.0%             |
| Always reuse (status quo) | none        | 100%       | 6.3%          | 100%     | 9.45%            |
| B3 GPTCache, $\tau=0.95$  | none        | 100%       | 6.3%          | 84.0%    | $\approx 9.45\%$ |
| B2 aggregate-only         | aggregate   | 97.8%      | 6.39% (37.2%) | 95.5%    | 9.37% (44.6%)    |
| B2’ SUPG-style, matched   | aggregate   | 95.5%      | 6.54%         | 95.5%    | 9.58%            |
| sIVM (this work)          | per stratum | 88.4%      | 3.28%         | 87.8%    | 5.47%            |

**Table 6: Model families on the lab column’s two edit pairs under the pinned procedure (presented-cells estimand,  $\delta=0.1$ ; the same seeded rows as the primary evaluation vector, one fresh draw per cell,  $T=0$ ). Cells give certified savings with the realized presented-cells error in parentheses, or “refused” with the oracle calls spent before the futility stop. The primary row is the main seed of Table 2 (whose savings are replication means). Gemini 3 Flash, priced at 23 times the output-token rate of the cheapest family here (gateway list prices, 29 September 2026), ran these two pairs only.**


| Family                                 | $n$   | floor | formatting-only ( $v1 \rightarrow v2$ ) |               |               | synonym rewording ( $v2 \rightarrow v3$ ) |               |                |
|----------------------------------------|-------|-------|-----------------------------------------|---------------|---------------|-------------------------------------------|---------------|----------------|
|                                        |       |       | flip                                    | $\alpha=0.1$  | $\alpha=0.2$  | flip                                      | $\alpha=0.1$  | $\alpha=0.2$   |
| DeepSeek V4 Flash (primary; main seed) | 2,000 | 5.0%  | 6.3%                                    | 88.4% (3.20%) | 88.4% (3.20%) | 8.9%                                      | refused (90)  | 82.3% (5.81%)  |
| Gemini 2.5 Flash-Lite                  | 2,000 | 4.9%  | 25.4%                                   | refused (90)  | refused (90)  | 31.5%                                     | 63.0% (0.84%) | 63.0% (0.84%)  |
| Gemini 3 Flash                         | 2,000 | 0.4%  | 4.5%                                    | 76.5% (0.88%) | 83.2% (1.05%) | 6.2%                                      | 69.5% (1.91%) | 85.6% (2.78%)  |
| GLM-4.7-Flash                          | 2,000 | 10.3% | 25.0%                                   | refused (90)  | refused (90)  | 31.3%                                     | refused (225) | 57.3% (9.64%)  |
| Qwen3.7-Flash                          | 2,000 | 5.5%  | 14.3%                                   | refused (765) | 61.8% (8.35%) | 27.2%                                     | 42.8% (3.67%) | 49.5% (4.25%)  |
| GPT-5 nano                             | 2,000 | 34.8% | 39.6%                                   | refused (90)  | refused (90)  | 39.7%                                     | refused (90)  | 41.3% (12.30%) |
| Gemini 3.8 Flash                       | 2,000 | 0.9%  | 0.9%                                    | 93.3% (0.65%) | 95.5% (0.70%) | 4.9%                                      | 90.5% (0.22%) | 90.5% (0.22%)  |

**Table 7: The remaining three edit pairs on further model families (same protocol as Table 6; floors are each family’s self-flip rate on that column, measured in the same run). The primary row is the main seed of Table 2; its widening-column floor is not shown, because the primary model’s self-flip rate under the production column’s first version was not measured (the lab column’s 5.0% is under the widened definition).**


| Family                      | SO scope widening (production column) |       |                                         |                | Djinni seniority tier (4-way select) |                                         |               |               |       |               |               |
|-----------------------------|---------------------------------------|-------|-----------------------------------------|----------------|--------------------------------------|-----------------------------------------|---------------|---------------|-------|---------------|---------------|
|                             | floor                                 | flip  | formatting-only ( $v1 \rightarrow v2$ ) |                | floor                                | criteria change ( $v2 \rightarrow v3$ ) |               |               |       |               |               |
|                             |                                       |       | $\alpha=0.1$                            | $\alpha=0.2$   |                                      | flip                                    | $\alpha=0.1$  | $\alpha=0.2$  | flip  | $\alpha=0.1$  | $\alpha=0.2$  |
| DeepSeek V4 Flash (primary) | –                                     | 9.4%  | 72.0% (4.50%)                           | 87.8% (5.33%)  | 6.5%                                 | 12.3%                                   | 19.1% (4.17%) | 52.3% (8.33%) | 12.4% | 18.9% (3.52%) | 62.7% (7.50%) |
| Qwen3.7-Flash               | 1.5%                                  | 17.2% | refused (135)                           | 79.8% (14.35%) | 2.8%                                 | 5.7%                                    | 57.3% (2.39%) | 78.1% (3.22%) | 6.1%  | 43.6% (3.33%) | 83.5% (4.80%) |
| Gemini 2.5 Flash-Lite       | 0.0%                                  | 8.6%  | refused (50)                            | 95.3% (8.22%)  | 0.2%                                 | 6.5%                                    | 23.3% (1.50%) | 86.5% (5.87%) | 5.1%  | 56.7% (2.73%) | 88.7% (4.57%) |
| GLM-4.7-Flash               | 9.4%                                  | 17.9% | 67.3% (6.03%)                           | 74.1% (6.29%)  | 4.4%                                 | 8.5%                                    | 36.5% (5.84%) | 65.4% (6.15%) | 9.3%  | 39.2% (1.26%) | 64.9% (3.94%) |
| Gemini 3.8 Flash            | 0.9%                                  | 4.6%  | 94.3% (4.35%)                           | 95.5% (4.35%)  | –                                    | –                                       | –             | –             | –     | –             | –             |

reuse-set guarantee applies only to certifications made in reuse-set mode. The strict reuse-set mode, whose deflated threshold shrinks with sampling depth, prices the stronger guarantee, and the price varies widely across the grid. On the headline edit at  $\alpha=0.2$  it is nearly free: 88.1% savings on the main seed against 88.4% in the default mode, still certifying in essentially every replication. One budget tighter it collapses, to 80.6% savings certified in 98.7% of replications, and on the unstable select column it never certifies reliably: no Djinni configuration at any budget certifies in more than 97.6% of replications. Guaranteeing the stale subset alone is therefore affordable exactly where the column is stable and the budget is loose, and unaffordable everywhere else.

*The  $\alpha$ -frontier and adaptive economics.* (Figure 2.) Savings collapse as  $\alpha$  approaches a stratum’s intrinsic rate, exactly as the power analysis predicts; below the floor the certifier refuses. The doubling schedule prices both directions: uncertifiable edits are refused after 90 oracle calls, whereas a fixed plan pays the full zero-flip minimum, and near-frontier certifications escalate sampling only while the data stays consistent with certifiability. A single refusal near the frontier is not evidence that the stratum exceeds the budget: the synonym edit at  $\alpha=0.1$  is refused on the main seed and certified in 81% of the sampling replications, a power miss the futility rule trades for its early stops.



**Figure 2: The  $\alpha$ -savings frontier under the pinned procedure: mean  $\pm$  sd over 1,000 sampling replications per point, generated from the result files. Savings collapse as  $\alpha$  approaches each column’s intrinsic flip rate.**



**Figure 3: Calibration power at  $\alpha=0.05$  (presented estimand, 2,000 trials per point): certification rises smoothly as  $p$  falls below  $\alpha$  and is small but nonzero just above it, where the realized population rate of a planted stratum often falls below  $\alpha$  and certifying it is correct; the unconditional rate of certificates whose realized error exceeds  $\alpha$  stays at or below 1.00% in every configuration against the nominal 5%. Power saturates once a stratum exceeds what the schedule can sample: the  $n_j=1800$  and  $n_j=5000$  curves nearly coincide because the six-look schedule caps both at the same sample size.**

*Deployment scale.* To test the claim that sample cost is table-size-independent, we materialized the Boolean column across the *entire* 89,184-row corpus: 78,956 billed cells in a single ledgered run costing \$9.31 at \$0.00012 per cell (the remaining 10,228 rows cost nothing, served by content-hash caching or already materialized by earlier passes), and 7.4 hours of model time reconstructed from the stored per-cell latencies at 32-way concurrency. We then certified the formatting edit against that materialized column with the real model as oracle, in three regimes over one sample path (Table 8): live in August 2026 under the Maurer–Pontil bound of the earlier version, which is the certificate the system applied; on the oracle draws that run stored, re-certified under the exact bound; and live again on 29 September 2026, 55 days later, under the exact bound, against the same cache.

*The August run.* Under Maurer–Pontil, at  $\alpha=0.2$  the procedure spent 225 oracle calls (certifying the cached-FALSE stratum at look 180 with 5 flips) to reuse 81,289 cells (91.1% fewer calls; audited error 3.16% over 2,686 audited reused rows), refusing the 7,715-row cached-TRUE stratum and so recomputing the 7,670 of its cells that the futility stop had not already drawn fresh; at  $\alpha=0.1$  it escalated to 765 calls (25 flips in 720) for 90.5% savings. The run applied the loose certificate, persisting a `reuse_certificate` row and stamping its identifier on 79,130 of the 81,289 reused cells, which is what the table then served; the 2,159 it did not stamp are the audited rows already recomputed for ground truth, because the apply step refuses to write over a freshly computed cell. A reviewer of the earlier version re-derived the loose certificate’s evidence (5 flips in 180 draws from the 81,469-row cached-FALSE stratum) under the Bardenet–Maillard bound and found 0.230, which does not clear the budget at the look it was issued; that is correct, and the exact bound reads 0.073 on the same evidence, which clears  $\alpha=0.2$  and  $\alpha=0.1$  alike. The tight certificate’s 25 flips in 720 read 0.055 under the exact bound, which clears  $\alpha=0.1$ , and 0.099 under Bardenet–Maillard, which clears it.

*The same draws under the exact bound.* Replayed over the stored August draws, the exact bound certifies the cached-FALSE stratum at look 45 (flips 1, bound 0.143), so the whole certification costs **90 oracle calls** and reuses **81,424 cells**: 91.3% fewer model calls (7,760 calls against 89,184), against a ceiling of 91.3% set by the cached-TRUE stratum the certifier never certifies, so the procedure captures essentially all the reuse available to capture and the residue is a property of the column. At  $\alpha=0.1$  it certifies at look 180 (5 flips, bound 0.073) for 225 calls, the count the Maurer–Pontil run needed for the loose budget, reusing 81,289 cells (91.1% savings). At  $\alpha=0.05$  the stored draws run out before the schedule’s deepest look (that look reaches 1,440 rows of the stratum, 699 of which hold no August draw), and drawing them later would mix two snapshots, so the August evidence is silent there. Realized error was audited against oracle-computed cells of that snapshot only (status done/cached; certificate copies are excluded by construction, as verbatim v1 values would register as non-flips): 89 flips in 2,821 audited reused rows, 3.15%, of which 1,812 rows belong to the released  $n=2,000$  evaluation vector and show 3.20% on their own; the released subset is a consistency check with Table 2, and the remainder is independent evidence. A separate verifier script recomputes the audit from the persisted reuse-set row identifiers. Priced at the ledger’s measured \$0.00012 per cell, the maintenance workload would cost \$0.91 and about 42 minutes of model time (per-call latencies summed, the same measure as the 7.4 hours; at 32-way concurrency the wall clock is a fraction of it), against the \$9.31 the corpus cost to materialize; we price the recompute leg without executing it.

*The September run: drift, measured.* Assumption 1 scopes a certificate to the model snapshot its oracle draws came from. We re-ran the certification live on 29 September 2026: same model identifier, same gateway, same prompts, same cache, same seeded path, with the oracle draws stored under a version of their own. The snapshot had moved. On the 720 cached-FALSE rows the August run had sampled, the edit’s flip rate against the cache read 3.5% in August and 6.9% in September; 36 rows flipped only in September against

**Table 8: Deployment-scale certification of the formatting edit on the 89,184-row column, one sample path, three regimes. “FALSE look / flips” is the look at which the cached-FALSE stratum was decided and the flips observed there; audited error is the flip rate among reused rows that hold an oracle cell of that snapshot.**

| Regime                                                                    | oracle calls                                                   | FALSE look / flips   | reused | savings | audited error (flips/rows) |
|---------------------------------------------------------------------------|----------------------------------------------------------------|----------------------|--------|---------|----------------------------|
| <i>August 2026 snapshot, oracle draws of the live run (version 2)</i>     |                                                                |                      |        |         |                            |
| Maurer–Pontil, $\alpha=0.2$ (live)                                        | 225                                                            | 180 / 5              | 81,289 | 91.1%   | 3.16% (85/2,686) (applied) |
| Maurer–Pontil, $\alpha=0.1$ (live)                                        | 765                                                            | 720 / 25             | 80,749 | 90.5%   | 3.03% (65/2,146)           |
| exact, $\alpha=0.2$ (replay)                                              | 90                                                             | 45 / 1               | 81,424 | 91.3%   | 3.15% (89/2,821)           |
| exact, $\alpha=0.1$ (replay)                                              | 225                                                            | 180 / 5              | 81,289 | 91.1%   | 3.16% (85/2,686)           |
| exact, $\alpha=0.05$                                                      | not reachable from the stored draws (699 further draws needed) |                      |        |         |                            |
| <i>September 2026 snapshot, oracle drawn live (version 4), same cache</i> |                                                                |                      |        |         |                            |
| exact, $\alpha=0.2$                                                       | 90                                                             | 45 / 2               | 81,424 | 91.3%   | 7.46% (104/1,395)          |
| exact, $\alpha=0.1$                                                       | 405                                                            | 360 / 20             | 81,109 | 90.9%   | 7.96% (86/1,080)           |
| exact, $\alpha=0.05$                                                      | 135                                                            | 90 / 6 (refused)     | 0      | 0.0%    | –                          |
| Maurer–Pontil, $\alpha=0.2$                                               | 225                                                            | 180 / 11             | 81,289 | 91.1%   | 7.54% (95/1,260)           |
| Maurer–Pontil, $\alpha=0.1$                                               | 1485                                                           | 1440 / 106 (refused) | 0      | 0.0%    | –                          |

11 only in August (McNemar  $p < 0.001$ ). A like-for-like comparison separates the model from the edit: the v1 prompt re-drawn in September disagreed with its own August draws, the cache, on 4.4% of those rows (Wilson [3.2, 6.2]), where the same prompt’s within-snapshot self-disagreement is 3.1% ([2.5, 3.7]); the v2 prompt re-drawn disagreed with its August draw on 6.5% ([4.9, 8.6]); and the edit’s flip rate with both sides drawn in September was 5.8% ([4.3, 7.8]), above its August value. The ledger shows the same change: 534 output tokens per cell against 290 in August, at 24 seconds per call against 9. On the 45 cached-TRUE rows the edit’s flip rate read 28.9% in August and 24.4% in September (8 rows flipped only in August, 6 only in September, McNemar  $p=0.79$ ), and the v2 prompt re-drawn disagreed with its August draw on 31.1% of them ([19.5, 45.7]; the stratum’s floor is 24.1%); too few rows to resolve a shift.

The certifier measures the sum of these effects against the cache it is asked to reuse, which is the quantity a user is exposed to, and it priced the drift as the theory says it should. At  $\alpha=0.2$  it certified the cached-FALSE stratum at look 45 (2 flips, against 1 in August): 90 calls in all, 81,424 cells reused (91.3%), 101 seconds of wall clock including the oracle calls; the cached-TRUE stratum was refused at look 45 with 11 flips. At  $\alpha=0.1$  it escalated to look 360 (20 flips, bound 0.091), one look deeper than the exact bound’s replay of the August draws: 405 calls, 81,109 cells (90.9%), 304 seconds. At  $\alpha=0.05$ , where the August draws are silent, the futility rule refused the stratum at look 90 (6 flips, 6.7%) after 135 calls. Under Maurer–Pontil the same September draws certify the loose budget after 225 calls and the tight budget refuses after 1485 calls. Audited against every September oracle cell, the reused rows flip at 7.46% (1,395 rows) under the loose certificate and 7.96% (1,080 rows) under the tight one: inside both budgets, and well above the August figure of 3.15%, which is the drift once more. The sample path needed 1,485 September oracle cells across the arms; 699 were already held from a pass earlier that day and 786 were requested from the runner, of which 20 were served by the content-hash cache and 766 were drawn fresh for \$0.15 (\$0.00020 per cell against \$0.00012 in August; the longer outputs are the difference). The drift measurement drew 752 more. A refusal here is the instrument reading correctly: a stratum whose flip rate against an aged cache exceeds the

budget is identified at the cost of one look, and its cells are recomputed rather than served.

The comparison to the  $n=2,000$  benchmark run shows that sampling cost did not scale with  $n$ : at  $\alpha=0.2$  the budget was drawn from the same distribution at both scales (90–180 calls, mean 97, against 2,000 rows; 90 against 89,184), because the doubling schedule stops on evidence about a stratum’s *rate*, which is size-independent, not on its size. The budget is a draw, since a schedule stops at the first look whose bound clears  $\alpha$ . At  $\alpha=0.2$  the benchmark’s main seed spent 90 calls, with a replication median of 90, and at  $\alpha=0.1$  the replications span 90–405 calls (mean 188, median 225). The scaling itself held: sampling cost is  $O(1)$  in  $n$  while the certifiable population is  $O(n)$ , so savings rose from 88.4% to 91.3% on the comparable single-run measurement (88.1% is the benchmark’s 1,000-replication mean, quoted elsewhere). Two reproducibility caveats. In MySQL the projection list changes which rows a seeded ORDER BY RAND returns, so the evaluation sample is defined by an exact query string, not by the seed alone. And the certifier of the earlier version shuffled each stratum in the order the cell query returned, which the query planner served through one index in August and another afterwards; the published run walked the content-hash order, the certifier now sorts by content hash before shuffling, and the replay reproduces the August run’s every count under the Maurer–Pontil arm (the released script asserts it).

*The stratification hierarchy.* Aggregate-only certification (one stratum) is valid for its aggregate estimand yet conceals 35.3% realized error in the reused cached-TRUE subgroup on the formatting edit at  $\alpha=0.1$  and 44.6% on the widening edit at  $\alpha=0.2$ : it buys its aggregate with the stratum whose floor is 24.1%, which per-value strata quarantine by construction. Refining further by embedding-interaction tertiles helps only when within-value risk is heterogeneous and otherwise costs power through the  $\delta/K$  split: on the headline configuration the value $\times$ embedding-tertile refinement certifies strictly less than value strata alone (82.7% vs 88.1% savings on SO formatting at  $\alpha=0.2$ ), because the split outweighs its discrimination gain. Across all 40 shared configurations of the released ablation the refinement is better in 6 and worse in

22, and the single win is a fraction of a point at a budget where both stratifiers certify essentially nothing. We report this as a negative result instead of tuning until it inverts: the refinement’s validity is unconditional (Assumption 2), its power is not, and on these columns within-value risk is homogeneous enough that the  $\delta/K$  split is never repaid.

*Minority-stratum anatomy.* The cached-TRUE stratum of the Boolean column flips 37.1% under a formatting edit; vote-of-3 on both sides reduces this to 24.7% and shows that 20.4% of cached TRUE values were single-draw noise; the flips that survive vote-3 are genuine boundary volatility. Quarantine-and-recompute is therefore the correct systemic response, and the actionable upstream knob is vote-based computation of boundary-stratum cells at write time.

*Calibration.* (Figure 3.) With planted flip rates across  $p \in [0, 0.25]$ , budgets  $\alpha \in \{0.05, 0.1, 0.2\}$  (including the one every headline result uses), stratum sizes  $\{600, 1800, 5000\}$ , both estimands, and 198 configurations (396,000 runs), the quantity Theorem 3.1 bounds is the rate of unsafe certificates: certifications of a stratum whose realized flip count exceeds  $\alpha n_j$  (in strict mode,  $\alpha(n_j - m_j)$ ). The worst such rate over all configurations is 1.00% (95% CI [0.65, 1.54]%;  $\alpha=0.05$ ,  $p=0.0525$ ,  $n_j=1,800$ , presented-cells mode) against the per-stratum  $\delta_j=5\%$  the study plants; the presented-cells rate, which leaves out the sampled flips, exceeds  $\alpha$  in at most 0.70% of runs. Of the configurations, 54 are *tight* nulls, planted just above the budget at 1.05 to 1.25 times  $\alpha$ , where the futility rule cannot dispose of the stratum at the first look and the sampling fraction climbs; the unsafe rate averages 0.35% over them, and the certification rate itself reaches 6.2%, most of it correct, because at  $p$  just above  $\alpha$  the realized population rate of a planted stratum falls below  $\alpha$  in a large share of trials and the exact bound certifies what is actually there. This is the regime that probes Assumption 3, and the bound holds its level in it with room to spare: the gap between 1.00% and 5% is the Bonferroni split across looks plus the discreteness of the test, and it is not budget that can be reclaimed without a sequential bound. The reuse-set rate evaluated on presented-mode certifications fluctuates past  $\alpha$  in rarely certifying frontier configurations (worst 100.0% of certified runs at  $p=0.0625$ ,  $\alpha=0.05$ , certified 0.80% of the time), as expected, since presented-mode does not certify that metric; the strict reuse-set mode holds it, with certification power rising smoothly as  $p$  falls below  $\alpha$  and with stratum size.

*Independence across rows.* Assumption 1 asks that oracle draws be independent across rows, and provider-side batching could in principle correlate concurrent requests [18, 47]. We tested the consequence the proof needs, non-interference, on 300 rows of the evaluation vector, drawing the formatting edit’s new version twice for every row on 29 September 2026: once with a single request in flight (78 minutes of wall clock) and once at the deployment run’s concurrency of 32 (4.6 minutes), against the same cache. Requests that errored or exceeded a 90-second ceiling are excluded from their arm: 9 sequential and 28 concurrent, where the ceiling binds more often. The falsification criterion was fixed in advance, at the 5% level throughout: a two-proportion test rejecting equal flip rates, a lag-one permutation test rejecting in the concurrent

arm, or a 95% CI on the concurrent arm’s block variance ratio that excludes 1. The rates agree: 8.9% (Wilson [6.2, 12.8]) sequential against 8.8% ([6.0, 12.8]) concurrent, two-proportion  $p=0.96$ , a test whose minimum detectable difference at 80% power is 6.8% absolute. The two arms disagree with each other on 9.5% of rows ([6.5, 13.7]), the same-prompt self-disagreement of the September snapshot (9.2% on 238 cached-FALSE rows, 12.0% on 25 cached-TRUE rows); the August floors belong to another snapshot and are not the comparison. In completion order, the lag-one autocorrelation of the flip indicator is -0.006 in the concurrent arm (permutation  $p=0.80$ ) and 0.070 in the sequential arm ( $p=0.22$ ). The variance of flip counts over consecutive blocks of 32 completions is 2.08 times its binomial value in the concurrent arm (8 blocks; 95% CI [0.91, 8.62], which covers 1; upper-tail chi-square  $p=0.04$ ) against 1.39 in the sequential arm (9 blocks;  $p=0.20$ ). No arm of the criterion is met, but the block test is the one we report as unresolved rather than passed: the one-sided tail is close to the level, and 8 blocks cannot see modest overdispersion. The rate test and the lag test find nothing at this concurrency through this gateway; a dependence at the block scale, and a rate difference smaller than 6.8%, are not excluded. Since the pinned schedule tests a stratum’s flip rate and the rates agree, the practical reading is that the certificates of the deployment run, which was issued at this concurrency, rest on draws whose rate the sequential arm reproduces.

*The bound.* Every result above spends its budget through the exact bound of Eq. (1), split across looks by Bonferroni. The bound, not the stratification or the schedule, is what sets the frontier, and we measured how much by swapping it inside the same pinned procedure over the same stored labels (Table 9; Appendix B lists every configuration). On a clean sample at  $n=90$  the exact bound reads 0.050 and its binomial limit, the Clopper–Pearson bound [12], 0.052; the Maurer–Pontil empirical-Bernstein bound [31] that an earlier version of this article certified with reads 0.144, the Bardenet–Maillard without-replacement bound [6, 39] 0.317, and a Waudby-Smith–Ramdas betting confidence sequence [46], which is anytime-valid and therefore spends the stratum’s budget without splitting it across the six looks, 0.065. End to end across the 20 configurations of Table 2 (main seed), the exact bound certifies 10 of them against 7 under Maurer–Pontil and 7 under Bardenet–Maillard, with mean savings 29.5% against 18.5% and 16.4%, on 7,458 oracle calls against 11,515 and 12,325. Relative to the exact bound, Maurer–Pontil loses 3 certificates and gains 0, and of the shared ones it reaches 1 at the same look and 6 at a later one; the widest single gain is SO scope widening at  $\alpha=0.1$ , from 18.1% savings for 1,485 calls to 72.0% for 405. The empirical-Bernstein bounds pay for a variance term a binary loss does not need, and the without-replacement form pays more, not less, because its constants are larger (an earlier preprint reported it as tighter; that replay used the Maurer–Pontil linear constant inside the Serfling form, and the published constant,  $\kappa = 7/3 + 3/\sqrt{2}$  [6], is what the numbers here use). The Clopper–Pearson limit certifies 9 configurations at mean savings 28.1% on 8,054 calls, the finite-population correction buying the exact bound its one further certificate. The betting sequence certifies 9 configurations at mean savings 29.1% on 7,616 calls; it beats the exact bound where the schedule’s Bonferroni split bites hardest (Djinni criteria change at

$\alpha=0.2$ : 76.3% against 62.7%), and loses to it elsewhere (certificates lost 1, gained 0). Its validity here rests on the i.i.d. construction of the hedged capital process, which the without-replacement draws do not satisfy exactly; a without-replacement betting sequence or a time-uniform boundary [20] would remove the split without that caveat and is the remaining power lever we see.

Validity comes before power, so all five arms were replayed through the calibration protocol (36,000 trials per arm, 18,000 of them at planted rates above the budget). No null configuration’s certification rate is demonstrably above the per-stratum  $\delta$  of 5% for any arm. The exact bound’s worst null certification rate is 6.3%. Most of those certificates are correct, because the realized count of a planted stratum falls below  $\alpha n_j$  in many trials; the certificates whose realized count exceeds the budget, the event of Theorem 3.1, are at most 1.00% of trials in any null configuration ( $\alpha=0.1$ ,  $p=0.105$ ,  $n_j=1,800$ ; Wilson [0.54, 1.83]%), and 0 of the 18 null configurations carry an interval above  $\delta$ . The Clopper–Pearson limit certifies 1.5% of null trials at worst (0.50% unsafe), the betting sequence 2.5% (1.20% unsafe), and Maurer–Pontil and Bardenet–Maillard 0.0% and 0.0% (0.00% and 0.00% unsafe), leaving almost all of their error budget unspent. A clean 1,800-cell stratum at  $\alpha=0.05$  costs the exact bound 90 draws (180 at  $n_j=5,000$  in the calibration study), the Clopper–Pearson limit 180, the sequence 180, Maurer–Pontil 360, and Bardenet–Maillard 720. The frontier in Figure 2 is therefore a property of the bound, and the exact bound is the pinned one because it is proven for exactly the sampling the procedure performs and, on every configuration measured here, is the tightest of the proven arms. Every number of the earlier version is reproducible by running the released harness with the Maurer–Pontil arm selected.

*Embedding-transfer operating points.* A measured  $\tau$  sweep over per-row similarities shows the guarantee-free transfer baseline collapses within a  $\approx 0.02$ -wide threshold window (formatting: 100% of rows reused at  $\tau=0.96$ , 10.7% at 0.98, none at 0.99; widening: 99.9% at 0.94, 8.1% at 0.96, none at 0.97), and, decisively, realized error among reused rows stays at the population flip rate at every operating point (5.6%–6.6% and 8.0%–9.8%); the threshold controls volume, never error, because cross-version similarity carries no per-cell flip signal (B1).

## 6 The prompt-edit benchmark

Semantic benchmarks exist, but they benchmark *engines* on a fixed task [25]; none exercises maintenance across prompt versions, which needs a different artifact: version pairs plus per-cell flip labels against the new version. Ours comprises: two redistributable corpora (Stack Overflow Developer Survey, 89,184 rows, ODbL 1.0/DbCL 1.0; Djinni English candidate profiles, 210,250 rows, MIT); three materialized columns of two types (two Boolean columns over structured facets, one of them the production column and one a lab replica seeded from its second version; one 4-way ordinal select over long free text); five realistic edit pairs spanning formatting-only, synonym rewording, scope widening, and criteria change; full ground-truth flip labels for  $n=2,000$  rows per version pair at pinned decode (28,000 labeled cells), plus self-flip replicate labels for both column types; and the deterministic harness (seeded sampling, adaptive certifier, baseline policies) that

reproduces every number in this paper. Table 10 lists the five pairs with their exact deltas. Ledgered model spend across the entire evaluation is \$12.01 (plus 99,957 cells from direct-call probes the ledger does not see), so the benchmark is cheap to extend to new models.

## 7 Related work

Table 11 positions sIVM against the lines of work it borrows from, is compared with, or is mistaken for, by the quantity each certifies and the axis of change it holds fixed.

*Guaranteed semantic query processing.* Semantic-operator systems bound the error of *one-shot* execution against a gold model: LOTUS’s optimized operators with accuracy targets [37], SUPG’s proxy-based selection with precision/recall guarantees [22] and its aggregation counterpart ABae, which stratifies by proxy before spending an expensive predicate [23], BARGAIN’s adaptive cascade thresholds [48], ThalamusDB’s approximate natural-language predicates with error bounds [21], and pipeline-level accuracy budgeting in Task Cascades [42]. All certify a *fixed* definition per query; none transfers evidence across a definition change. We reuse their statistical framing (a labeled sample, a bound, a threshold) in a maintenance loop those systems do not model, with a different bound: because the loss is binary and the sample is drawn without replacement from a frozen stratum, the exact finite-population bound, the hypergeometric form of Clopper and Pearson’s binomial limits [12], is available and dominates the empirical-Bernstein bounds those systems and an earlier version of this work relied on (§5); time-uniform confidence sequences [20, 46] are the alternative that would remove the split across looks. The certificate itself is distribution-free risk control in the sense of risk-controlling prediction sets [7] and Learn then Test [2]: an upper confidence bound on a risk that must clear a user-set  $\alpha$  with probability  $1 - \delta$ , here evaluated per stratum and per look with the multiple-testing accounting those frameworks make explicit. The declarative LLM data systems that optimize cost and quality without a statistical target, Palimpsest’s AI-powered analytics [28], DocETL’s agentic pipeline rewriting [40], Evaporate’s structured views over data lakes [3], and UQE’s sampling-driven queries over unstructured collections [13], compute the same kind of prompt-defined per-record outputs, and none certifies which of its outputs survive a redefinition.

*Verified caching.* vCache certifies that a *new query* may reuse a cached answer under one fixed definition, with learned per-entry thresholds and user error rates [38]; Evergreen spends confidence sequences and early stopping to cut the cost of *verifying* semantic aggregates, not to decide reuse [26]; and FreshCache approves a cache hit only when a learned staleness model puts the probability that the entry is stale below a per-tier error budget [29]. FreshCache is the closest in shape, since a per-partition budget gating reuse is exactly our skeleton, and it differs on the two axes that decide the guarantee: its risk is change-of-world (open-web evidence aging under fixed query semantics) where ours is change-of-definition, and its per-tier quantity is a fitted model’s point estimate where ours is a finite-sample upper confidence bound that

**Table 9: Five upper confidence bounds inside the same pinned procedure on the same stored labels: the bound on a clean sample of 90 draws and on the deployment evidence of 5 flips in 180 draws (the four fixed-sample bounds at the pinned per-look level,  $\delta/12$ ; the anytime-valid betting sequence at the unsplit per-stratum level,  $\delta/2$ ); the certifying configurations, mean savings, and oracle calls over the Table 2 grid (main seed); and the worst rate of unsafe certificates, the event of Theorem 3.1, at planted nulls in the calibration replay (per-stratum  $\delta=0.05$ ).**

| Bound                            | upper bound on |                  | Table 2 grid (20 configurations) |              |              | unsafe at nulls<br>worst |
|----------------------------------|----------------|------------------|----------------------------------|--------------|--------------|--------------------------|
|                                  | 0/90, $N=1800$ | 5/180, $N=81469$ | certifying                       | mean savings | oracle calls |                          |
| Exact finite-population (pinned) | 0.050          | 0.073            | 10                               | 29.5%        | 7,458        | 1.00%                    |
| Clopper–Pearson (binomial)       | 0.052          | 0.073            | 9                                | 28.1%        | 8,054        | 0.50%                    |
| Betting confidence sequence      | 0.065          | –                | 9                                | 29.1%        | 7,616        | 1.20%                    |
| Maurer–Pontil (first submission) | 0.144          | 0.140            | 7                                | 18.5%        | 11,515       | 0.00%                    |
| Bardenet–Maillard (WoR)          | 0.317          | 0.230            | 7                                | 16.4%        | 12,325       | 0.00%                    |

The betting sequence has no 5/180 entry because its bound is a function of the whole draw sequence, not of  $(k, n)$  alone.

**Table 10: The five versioned prompt-edit pairs. Deltas are transcribed from the stored prompt templates; bound fields are the row attributes the template interpolates. Flip rates are single-fresh-draw labels against the cached version on the  $n=2,000$  evaluation vector (Table 2); the SO synonym pair is measured against the formatting version’s cells (v2→v3), as the benchmark defines it.**

| Pair                           | Column (corpus; contract; bound fields)                                                            | Exact delta                                                                                                                                                                                                                                                                                                                                                                           | flip   |
|--------------------------------|----------------------------------------------------------------------------------------------------|---------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------|--------|
| SO formatting-only (v1→v2)     | Data-platform specialist, lab replica (SO; Boolean; role, databases, platforms, tools, other tech) | Two paragraph breaks inserted after the question and before the “Judge from” sentence; “OR heavy analytics engineering” lowercased to “or heavy analytics engineering”; no content word changed                                                                                                                                                                                       | 6.3%   |
| SO synonym rewording (v2→v3)   | same                                                                                               | “this person” → “this individual”; “works primarily on” → “chiefly builds and maintains”; “heavy analytics engineering” → “intensive analytics engineering”; “Judge from role ...and other tech” → “Decide using their role ...and other technologies”; “Working with multiple ...counts as evidence” → “Experience with several ...is supporting evidence”                           | 8.90%  |
| SO scope widening (v1→v2)      | Data-platform specialist, production column (SO; Boolean; same fields)                             | Definition extended from “warehouses, lakes, pipelines, distributed data stores, or ML data tooling” by “OR heavy analytics engineering”; the exclusion “Merely using one relational database for app work does not qualify” replaced by the inclusion “Working with multiple analytical/distributed databases (BigQuery, Snowflake, Spark, Kafka, Elasticsearch) counts as evidence” | 9.45%  |
| Djinni formatting-only (v1→v2) | Seniority tier, lab (Djinni; 4-way ordinal select; position, years of experience, CV text)         | The one-line template broken into a question line, three labeled input lines (Position, Years of experience, CV), and the instruction sentence, with no word changed                                                                                                                                                                                                                  | 12.30% |
| Djinni criteria change (v2→v3) | same                                                                                               | Sentence appended: “Evidence of team or people leadership (leading engineers, owning hiring, running a team) moves the candidate up one tier”                                                                                                                                                                                                                                         | 12.45% |

either clears  $\alpha$  or refuses. Our B1 baseline shows why these mechanisms do not port: across a definition change the query-to-entry similarity that such caches threshold carries no per-cell flip signal (AUC 0.454 and 0.499). sIVM certifies change-of-*definition* transfer, the axis those systems hold fixed; the two compose (a certificate is scoped to a snapshot, and change-of-world machinery decides when the snapshot itself expires).

*Incremental and adaptive view maintenance.* Sampling a stale materialized view to bound the error of answers served from it [24] is the closest classical relative of what we do, and the difference is what goes stale: there the data drifts under a fixed view definition, here the definition itself changes over fixed data, which is why their staleness estimators condition on updates and ours condition

on a prompt delta, and the look schedule itself is online aggregation’s template [19], a running estimate whose interval tightens as sampling proceeds, repurposed to certify a stratum rather than to answer an aggregate. Classical incremental view maintenance (IVM) and view adaptation after redefinition [8, 17] assume deterministic operators, where a rewrite either preserves a subexpression or does not. Stochastic LLM cells break the premise: our measurements show that even a semantically null rewrite flips cells at a rate above the column’s noise floor (6.3% against 5.0%), so “preserved” is not a property a rewrite can have here, only a rate one can bound. Streaming semantic systems recompile operators mid-stream but recompute outputs instead of certifying reuse of old ones [10, 11], and streaming cascades certify on the *data-arrival*

**Table 11: What neighboring systems certify, and across what. sIVM is the only row whose certified quantity is the reuse of cached outputs across a change of the prompt definition.**

| System or line                                                                | Certified quantity                                                                     | Held fixed                      | Changes                        | Guarantee form                                        |
|-------------------------------------------------------------------------------|----------------------------------------------------------------------------------------|---------------------------------|--------------------------------|-------------------------------------------------------|
| SUPG, ABae, BARGAIN, LOTUS, ThalamusDB, Task Cascades [21–23, 37, 42, 48]     | precision, recall, or accuracy of one query’s answer against an oracle                 | the query definition            | nothing (one-shot execution)   | finite-sample bound per query                         |
| vCache [38]                                                                   | correctness of serving a cached answer to a <i>new</i> query                           | the answer semantics            | the incoming query             | learned per-entry threshold at a user error rate      |
| FreshCache, Evergreen [26, 29]                                                | probability that a cached entry is stale (world drift); cost of verifying aggregates   | the query definition            | the world the answer describes | fitted staleness model per tier; confidence sequences |
| Stale-view cleaning, online aggregation [19, 24]                              | error of an answer served from a stale view; a running aggregate                       | the view definition             | the base data                  | sampling bound on the answer                          |
| Classical IVM, view adaptation [8, 17]                                        | exact equivalence of a rewritten view                                                  | deterministic operators         | the view definition            | algebraic (no sampling)                               |
| Step-level reuse [34]                                                         | per-step equivalence under a perturbed request                                         | the model                       | the request                    | deterministic per-step check                          |
| Guarantee-free semantic caches [5] (B3)                                       | nothing: a hit is a similarity above a threshold                                       | nothing                         | the incoming query             | none                                                  |
| Per-output verification: assertions, regression suites, vote-of- $k$ [14, 41] | correctness of one output against a checker or a majority                              | the prompt                      | the output under test          | per-output check, no population bound                 |
| sIVM (this work)                                                              | false-reuse rate of cached stochastic cells kept under the new definition, per stratum | the model snapshot and the data | the prompt definition          | exact finite-sample bound per stratum and look        |

axis, where new tuples meet a fixed operator, and leave the definition axis open [27]. Given that distinction, sIVM is, to our knowledge, the first maintenance procedure whose unit of transfer is a *statistically certified set of stochastic cells*.

The nearest counterexample shape is step-level reuse with lightweight verification and selective patching across perturbed requests [34], which does reuse computation when a request changes. It differs on exactly the property this paper is about: its verification is a deterministic per-step check, so it carries no error budget and makes no population statement about the set it reuses. In-database prompt management [30] and versioned, provenance-tracked prompt views [9] supply the substrate this work assumes, not a competing mechanism; we take the framing of prompts as versioned views from that line and claim only the certification machinery built on top of it, whose reuse is of *outputs*, not of key-value (KV) attention state.

*Prompt sensitivity and regression.* Prompt-delta analysis for assertion synthesis [41] anticipates our edit taxonomy but targets pipeline quality checks, not reuse certification; interactive LLM regression tools frame migration pain without guarantees [14]. Our finding that surface edit class can fail to predict flip rate (a synonym paraphrase flipping more cells than a formatting edit; a formatting edit indistinguishable from a criteria change) shows why measurement, not taxonomy, must carry the decision.

## 8 Limitations and outlook

*Which budgets are usable, and for what.* Most of the savings in Table 2 sit at  $\alpha=0.2$ , which allows up to one cell in five of a certified stratum to flip on recomputation (with probability  $1 - \delta$  over

the run, in the realized terms of Theorem 3.1). That is a defensible budget where a person reads the column row by row before acting on it (a recruiter screen, a lead list, an enrichment sheet a reviewer walks through), where the cost of a stale cell is a wasted glance, and where the alternative on offer is today’s status quo, which reuses everything at the population flip rate with no bound at all. It is not a defensible budget for a column that feeds an aggregate or a downstream filter; there the strict reuse-set mode and a budget at or below 0.05 are the right setting, and §5 shows what they cost: on the stable Boolean column  $\alpha=0.1$  certifies 82.6% mean savings and  $\alpha=0.05$  certifies in 74.4% of replications at the deepest look, while on the less stable select column tight budgets are refused because the model’s own floor sits above them. At deployment scale the sample a tight budget needs is still a constant, 405 calls at  $\alpha=0.1$  against 89,184 cells, so tight budgets are affordable in calls; what they cannot buy is certification of a stratum whose flip rate exceeds them, which is the correct outcome. A practitioner should therefore read the frontier of Figure 2 against the column’s floor and choose the budget the downstream use tolerates, not the one that maximizes savings.

*Statistical scope.* Bonferroni across strata and looks is sound but conservative, and §5 prices it: the exact bound leaves a gap between its measured unsafe-certificate rate and the nominal  $\delta$ , and a betting confidence sequence, which spends no split across looks, beats it on the configurations where that split bites hardest and loses elsewhere. A without-replacement betting sequence or a time-uniform boundary [20] would combine the two, and e-value selection (the e-value Benjamini–Hochberg procedure, e-BH) [45] composes with either; neither is needed for the guarantees stated



here. Our stratifiers are first instances; the lexical variant is weak on broad edits, and the embedding variant’s signal exists only for semantic edits, so learned interaction models are an open lever with guaranteed validity (Assumption 2).

*Equivalence and columns.* In the main evaluation equivalence is exact match, and the columns are two Boolean columns and one four-way ordinal select; numeric and free-text columns, which the introduction’s product examples also serve, are not evaluated there. A free-text column needs a judge, and the judge’s own error enters the budget. Let the judge miss a fraction  $q$  of true flips (call a fresh draw equivalent to the cache when a reader would not). Calibrate  $q$  on human-labeled pairs from the same (edit, column) with an upper confidence bound  $\hat{q}$  at level  $\delta_{\text{cal}}$ , and certify at  $\alpha - \hat{q}$ : the flip rate the judge reports is at least the true rate minus  $q$ , so a certified stratum’s true rate is at most  $\alpha$  whenever the calibration bound and the certificate both hold, which they do with probability at least  $1 - \delta - \delta_{\text{cal}}$ . Errors in the other direction (calling equivalent cells different) only raise the reported rate and cost power. The calibration is per (edit, column), since a judge’s miss rate is a property of the outputs it compares; the machinery of this article is otherwise unchanged, but its power under a noisy judge is an open measurement. The main evaluation’s five author-written edit pairs, one per class per column, are enough to show that surface class can fail to predict flip rate, not to estimate how often it does, and the edit-class findings should be read at that strength.

*Models and snapshots.* One budget-class reasoning model carries the full experimental program; 6 further families were run on the lab column’s two pairs and 3 of them on all five (§5), where certification reproduced on 6 of them and the same edit’s flip rate ranged widely, so the quantitative frontier (where savings sit for a given  $\alpha$ , and which edits are benign at all) is specific to the (edit, model) pair, even though the qualitative structure, a measurable floor that bounds certifiability and a certifier whose refusals track the measured flip rate, held on every family. Certificates are scoped to the model snapshot their oracle draws came from. §5 shows what that scope is worth: the same model identifier through the same gateway had drifted 55 days after the deployment run; the certifier measured the drift as a higher flip rate against the old cache, re-certified the two budgets the August run had certified (one of them a look deeper than the exact bound’s replay of the August draws), and refused the tightest budget, on which the August evidence is silent. Detecting when a snapshot has moved without a full re-certification, i.e., composing sIVM with change-of-world invalidation [26, 29], is the natural next step.

*Economics.* The economic frame that matters in production is time-to-consistency under provider rate limits, where certified maintenance samples 5%–74% of cells across the reliably certifying configurations of Table 2 (5% on the headline formatting edit at  $\alpha=0.2$ ) instead of recomputing 100%, a wall-clock and rate-limit gain that grows with table size, because sample sizes do not; the live re-certification of §5 measured that decision at 101 seconds of wall clock for the loose budget. Two things this article does not measure. The certificate surface in the UI was not studied with users, so whether a person acts differently on a column that shows

one is open. And the cost figures come from one ledgered materialization and one priced maintenance replay at the gateway’s list rates; they show the scale of the saving, not a cost model.

## 9 Conclusion

Prompt-defined columns are materialized views whose *definition* changes far more often than their data, and the systems serving them today choose between recomputing everything and reusing everything. sIVM shows the middle is reachable and cheap: freeze a stratification, sample on a doubling schedule, bound each stratum’s flip rate with an exact finite-population bound, reuse only what clears the budget. At deployment scale that cost 90 oracle calls to certify reuse of 81,424 of 89,184 cells, and the sample did not have to grow with the table, because the schedule stops on evidence about a *rate*, which is size-independent.

The result we did not anticipate is that the certifier doubles as a measuring instrument for the column it maintains. A stratum refused at every budget is a slice of the table on which the model cannot reproduce its own answers, identified by a futility stop at the cost of one look; a flip rate that has risen on the same rows against the same cache 55 days later is a model that has moved under it, and the certifier reports the movement in the currency the user pays, oracle calls per certified budget. Certified maintenance is therefore useful even where it certifies nothing: the refusal identifies, cheaply, reuse error that would otherwise go unmeasured.

## Availability and reproducibility

The system, the certifier, every experiment script, and every result artifact are public in a single repository, <https://github.com/arjunlohan/project-lore>, under a source-available license that permits research and noncommercial use. Released: the two corpora (already public and redistributable under ODbL 1.0/DbCL 1.0 and MIT), the five versioned edit pairs, the per-cell ground-truth flip labels, the certifier, the adapter conformance suite, and every script and result artifact that produces a number in this paper. Both corpora are used as distributed by their maintainers under those licenses; the Djinni corpus is released anonymized at source, and this work collects no new human-subject data, attempts no re-identification, and releases only per-cell labels and aggregate statistics. Nothing in the evaluation path is withheld. Anyone reproducing this needs a model endpoint of their own; the labels are released so the statistical results can be reproduced without one. Several numbers here (the per-stratum floors, the ledger economics, the materialization wall clock, the certificate’s stamp count) are computed from the running system’s database, not from a result file, so the generator dumps every one of those database reads to an artifact alongside the others; the release is therefore self-contained, with no dependence on our instance. Every quantitative claim in this paper’s prose, tables, and figures is emitted by a single generator (`scripts/experiments/gen-paper-assets.ts`) from those artifacts and from the system’s cost ledger, and cited through macros, never typed inline, so re-running an experiment regenerates the sentences that cite it. A checker (`scripts/check-paper-numbers.ts`) enforces this mechanically and fails on nine classes of defect, listed in Appendix A; it reports zero of all nine for this draft. Three reproducibility

hazards we hit are documented in the harness and in §5: seeded ORDER BY RAND in MySQL depends on the projection list, so evaluation samples are pinned by exact query string; the certifier’s sample path depended on the storage engine’s row order until it was pinned to a canonical order; and naive linear-congruential pseudorandom number generators written with JavaScript number multiplication silently lose low bits, so all sampling uses exact 32-bit arithmetic. Relative to the Zenodo preprint of this article, the certifier’s default bound changed from Maurer–Pontil to the exact finite-population bound and every number was regenerated; the earlier numbers are reproducible from the same release with the Maurer–Pontil arm selected.

## A The mechanical guards

scripts/check-paper-numbers.ts fails on any of nine conditions: (1) a hand-typed quantity in the prose outside an enumerated allow-list (user parameters such as  $\alpha$  and  $\delta$ , protocol constants such as the schedule looks, and generic thresholds quoted as such); (2) a quantity hardcoded inside the generator instead of derived from an artifact; (3) a macro defined but never cited, since an uncited macro can drift undetected; (4) a macro whose trailing space is swallowed into the next word, which is invisible in the source; (5) a sentence that asserts two quantities differ while citing two macros whose values are equal, which is how an earlier draft printed “different looks” beside a number repeated twice; (6) a spelled-out multiplier (“twice”, “an order of magnitude”) that the cited macros’ ratio does not support, which is how a rate 1.26 times a floor was printed as double it; (7) a containment claim (“inside”, “at most”) that the cited range and scalar macros do not support; (8) an absolute claim (“never”, “every”, “anywhere”) appearing in the abstract or the contributions list without an entry in an enumerated allow-list, the one class of defect the numeric guards are blind to by construction; and (9) a scope claim (“on the same column”) whose cited macros come from different columns, checked against a provenance file the generator emits from the same pair definitions the experiments run on. Each guard exists because the corresponding defect reached a compiled PDF at least once, several of them in a draft whose numbers were already machine-generated: generating the quantities is necessary but not sufficient, because the sentences that *relate* quantities go stale independently of the quantities themselves.

## B The bound ablation, configuration by configuration

Table 12 gives, for every configuration of Table 2 and each of the five bounds, the certified savings of the main seed and the oracle calls spent. Relative to the exact bound, the Clopper–Pearson bound certifies all but one of its configurations (lost 1, gained 0), reaching 7 of them at the same look, 0 earlier, and 2 later, since it ignores the finite-population correction; Maurer–Pontil reaches 1 at the same look, 6 later, and 0 earlier; Bardenet–Maillard loses 3 certificates, gains 0, and reaches 0 at the same look, 0 earlier, and 7 later; the betting sequence reaches 7 at the same look, 1 earlier, and 1 later.

## References

- [1] Airtable, Inc. 2026. Using Airtable AI in Fields. Airtable Support, <https://support.airtable.com/articles/8052242094-using-airtable-ai-in-fields>. Accessed 29 September 2026.
- [2] Anastasios N. Angelopoulos, Stephen Bates, Emmanuel J. Candès, Michael I. Jordan, and Lihua Lei. 2025. Learn then Test: Calibrating Predictive Algorithms to Achieve Risk Control. *Ann. Appl. Stat.* 19, 2 (2025), 1641–1662. <https://doi.org/10.1214/24-AOAS1998> arXiv:2110.01052.
- [3] Simran Arora, Brandon Yang, Sabri Eyuboglu, Avanika Narayan, Andrew Hojell, Immanuel Trummer, and Christopher Ré. 2023. Language Models Enable Simple Systems for Generating Structured Views of Heterogeneous Data Lakes. *Proc. VLDB Endow.* 17, 2 (2023), 92–105. <https://doi.org/10.14778/3626292.3626294>
- [4] Berk Atıl, Sarp Aykent, Alexa Chittams, Lisheng Fu, Rebecca J. Passonneau, Evan Radcliffe, Guru Rajan Rajagopal, Adam Sloan, Tomasz Tudrej, Ferhan Ture, Zhe Wu, Lixinyu Xu, and Breck Baldwin. 2024. Non-Determinism of “Deterministic” LLM Settings. arXiv:2408.04667
- [5] Fu Bang. 2023. GPTCache: An Open-Source Semantic Cache for LLM Applications. In *Proc. Workshop for Natural Language Processing Open Source Software (NLP-OSS) at EMNLP*.
- [6] Rémi Bardenet and Odalric-Ambrym Maillard. 2015. Concentration Inequalities for Sampling without Replacement. *Bernoulli* 21, 3 (2015), 1361–1385. <https://doi.org/10.3150/14-BEJ605>
- [7] Stephen Bates, Anastasios Angelopoulos, Lihua Lei, Jitendra Malik, and Michael I. Jordan. 2021. Distribution-Free, Risk-Controlling Prediction Sets. *J. ACM* 68, 6 (2021), 43:1–43:34. <https://doi.org/10.1145/3478535>
- [8] Mihai Budiş, Frank McSherry, Leonid Ryzhyk, and Val Tannen. 2023. DBSP: Automatic Incremental View Maintenance for Rich Query Languages. *Proc. VLDB Endow.* 16, 7 (2023), 1601–1614. <https://doi.org/10.14778/3587136.3587137>
- [9] Uğur Çetintemel, Shu Chen, Alexander W. Lee, Deepti Raghavan, Duo Lu, and Andrew Crotty. 2026. Making Prompts First-Class Citizens for Adaptive LLM Pipelines. In *Conf. on Innovative Data Systems Research (CIDR)*. arXiv:2508.05012.
- [10] Shu Chen, Junhan Liu, Deepti Raghavan, and Uğur Çetintemel. 2026. VecFlow: Long-Horizon Semantic Processing over Data and Event Streams with LLMs. arXiv:2604.03855
- [11] Shu Chen, Deepti Raghavan, and Uğur Çetintemel. 2025. Continuous Prompts: LLM-Augmented Pipeline Processing over Unstructured Streams. arXiv:2512.03389
- [12] C. J. Clopper and E. S. Pearson. 1934. The Use of Confidence or Fiducial Limits Illustrated in the Case of the Binomial. *Biometrika* 26, 4 (1934), 404–413. <https://doi.org/10.1093/biomet/26.4.404>
- [13] Hanjun Dai, Bethany Yixin Wang, Xingchen Wan, Bo Dai, Sherry Yang, Azade Nova, Pengcheng Yin, Phitchaya Mangpo Phothilimthana, Charles Sutton, and Dale Schuurmans. 2024. UQE: A Query Engine for Unstructured Databases. In *Advances in Neural Information Processing Systems (NeurIPS)*. arXiv:2407.09522.
- [14] Tanay Dixit, Daniel Lee, Sally Fang, Sai Sree Harsha, Anirudh Sureshan, Akash Maharaj, and Yunyao Li. 2024. RETAIN: Interactive Tool for Regression Testing Guided LLM Migration. In *Proc. Conf. on Empirical Methods in Natural Language Processing (EMNLP): System Demonstrations*. 301–310. arXiv:2409.03928.
- [15] Fabrizio Giliardi, Meysam Alizadeh, and Maël Kubli. 2023. ChatGPT Outperforms Crowd Workers for Text-Annotation Tasks. *Proc. Natl. Acad. Sci. USA* 120, 30 (2023), e2305016120. <https://doi.org/10.1073/pnas.2305016120>
- [16] Google LLC. 2026. Use the AI Function in Google Sheets. Google Docs Editors Help, <https://support.google.com/docs/answer/15877199>. Accessed 29 September 2026.
- [17] Ashish Gupta, Inderpal Singh Mumick, and Kenneth A. Ross. 1995. Adapting Materialized Views after Redefinitions. In *Proc. ACM SIGMOD Int. Conf. on Management of Data*. 211–222. <https://doi.org/10.1145/223784.223817>
- [18] Horace He and Thinking Machines Lab. 2025. Defeating Nondeterminism in LLM Inference. Thinking Machines Lab: Connectionism, technical report. <https://doi.org/10.64434/tml.20250910> Accessed 15 August 2026.
- [19] Joseph M. Hellerstein, Peter J. Haas, and Helen J. Wang. 1997. Online Aggregation. In *Proc. ACM SIGMOD Int. Conf. on Management of Data*. 171–182. <https://doi.org/10.1145/253260.253291>
- [20] Steven R. Howard, Aaditya Ramdas, Jon McAuliffe, and Jasjeet Sekhon. 2021. Time-Uniform, Nonparametric, Nonasymptotic Confidence Sequences. *Ann. Statist.* 49, 2 (2021), 1055–1080. <https://doi.org/10.1214/20-AOS1991>
- [21] Saehan Jo and Immanuel Trummer. 2024. ThalamusDB: Approximate Query Processing on Multi-Modal Data. *Proc. ACM Manag. Data* 2, 3 (2024), 1–26. <https://doi.org/10.1145/3654989>
- [22] Daniel Kang, Edward Gan, Peter Bailis, Tatsunori Hashimoto, and Matei Zaharia. 2020. Approximate Selection with Guarantees Using Proxies. *Proc. VLDB Endow.* 13, 12 (2020), 1990–2003. <https://doi.org/10.14778/3407790.3407804>
- [23] Daniel Kang, John Guibas, Peter Bailis, Tatsunori Hashimoto, Yi Sun, and Matei Zaharia. 2021. Accelerating Approximate Aggregation Queries with Expensive Predicates. *Proc. VLDB Endow.* 14, 11 (2021), 2341–2354. <https://doi.org/10.14778/3476249.3476285>

**Table 12: Certified savings of the main seed (oracle calls in parentheses) for every configuration and bound; “refused” gives the calls spent before the stop.**

| Edit                   | $\alpha$ | exact         | Clopper-Pearson | betting       | Maurer-Pontil  | Bardenet-Maillard |
|------------------------|----------|---------------|-----------------|---------------|----------------|-------------------|
| SO formatting-only     | 0.02     | refused (225) | refused (225)   | refused (225) | refused (225)  | refused (225)     |
| SO formatting-only     | 0.05     | 18.7% (1485)  | 18.7% (1485)    | 18.7% (1485)  | refused (1485) | refused (1485)    |
| SO formatting-only     | 0.10     | 88.4% (90)    | 86.2% (135)     | 86.2% (135)   | 54.7% (765)    | 54.7% (765)       |
| SO formatting-only     | 0.20     | 88.4% (90)    | 88.4% (90)      | 88.4% (90)    | 86.2% (135)    | 72.7% (405)       |
| SO synonym rewording   | 0.02     | refused (90)  | refused (90)    | refused (90)  | refused (90)   | refused (90)      |
| SO synonym rewording   | 0.05     | refused (90)  | refused (90)    | refused (90)  | refused (90)   | refused (90)      |
| SO synonym rewording   | 0.10     | refused (90)  | refused (90)    | refused (90)  | refused (90)   | refused (90)      |
| SO synonym rewording   | 0.20     | 82.3% (225)   | 82.3% (225)     | 82.3% (225)   | 82.3% (225)    | 73.3% (405)       |
| SO scope widening      | 0.02     | refused (90)  | refused (90)    | refused (90)  | refused (90)   | refused (90)      |
| SO scope widening      | 0.05     | refused (765) | refused (765)   | refused (765) | refused (765)  | refused (765)     |
| SO scope widening      | 0.10     | 72.0% (405)   | 72.0% (405)     | 72.0% (405)   | 18.1% (1485)   | 18.1% (1485)      |
| SO scope widening      | 0.20     | 87.8% (90)    | 87.8% (90)      | 87.8% (90)    | 81.0% (225)    | 72.0% (405)       |
| Djinni formatting-only | 0.02     | refused (180) | refused (180)   | refused (180) | refused (180)  | refused (180)     |
| Djinni formatting-only | 0.05     | refused (315) | refused (315)   | refused (315) | refused (315)  | refused (315)     |
| Djinni formatting-only | 0.10     | 19.1% (495)   | refused (878)   | refused (878) | refused (878)  | refused (878)     |
| Djinni formatting-only | 0.20     | 52.3% (953)   | 52.3% (953)     | 52.3% (953)   | 19.1% (1617)   | 19.1% (1617)      |
| Djinni criteria change | 0.02     | refused (180) | refused (180)   | refused (180) | refused (180)  | refused (180)     |
| Djinni criteria change | 0.05     | refused (315) | refused (315)   | refused (315) | refused (315)  | refused (315)     |
| Djinni criteria change | 0.10     | 18.9% (540)   | 18.9% (540)     | 18.9% (540)   | refused (918)  | refused (918)     |
| Djinni criteria change | 0.20     | 62.7% (745)   | 54.4% (913)     | 76.3% (475)   | 27.9% (1442)   | 18.9% (1622)      |

- [24] Sanjay Krishnan, Jiannan Wang, Michael J. Franklin, Ken Goldberg, and Tim Kraska. 2015. Stale View Cleaning: Getting Fresh Answers from Stale Materialized Views. *Proc. VLDB Endow.* 8, 12 (2015), 1370–1381. <https://doi.org/10.14778/2824032.2824037>
- [25] Jiale Lao et al. 2026. SemBench: A Benchmark for Semantic Query Processing Engines. arXiv:2511.01716 To appear, VLDB 2026.
- [26] Alexander W. Lee, Benjamin Han, Shayak Sen, Sam Yeom, Ugur Cetintemel, and Anupam Datta. 2026. Evergreen: Efficient Claim Verification for Semantic Aggregates. arXiv:2604.26180
- [27] Paweł Liskowski and Kyle Schmaus. 2026. Streaming Model Cascades for Semantic SQL. arXiv:2604.00660
- [28] Chunwei Liu, Matthew Russo, Michael Cafarella, Lei Cao, Peter Baile Chen, Zui Chen, Michael Franklin, Tim Kraska, Samuel Madden, Rana Shahout, and Gerardo Vitagliano. 2025. Palimpsest: Optimizing AI-Powered Analytics with Declarative Query Processing. In *Conf. on Innovative Data Systems Research (CIDR)*.
- [29] Muhammad Mansoor, Tahir Ahmad, and Yeo-Chan Yoon. 2026. Risk-Constrained Freshness-Aware Semantic Caching for Open-Web Retrieval-Augmented LLMs. arXiv:2607.04281
- [30] Denis Mayr Lima Martins and Gottfried Vossen. 2026. Prompt as a Data Type: In-Database LLM Prompt Management and Rewriting. arXiv:2607.21756
- [31] Andreas Maurer and Massimiliano Pontil. 2009. Empirical Bernstein Bounds and Sample-Variance Penalization. In *Proc. Conf. on Learning Theory (COLT)*. Online proceedings, paper 012; no DOI or page numbers are assigned; also arXiv:0907.3740.
- [32] Microsoft Corporation. 2026. COPILOT Function. Microsoft Support, <https://support.microsoft.com/en-us/excel/functions/copilot-function>. Preview feature, withdrawn 14 September 2026; accessed 29 September 2026.
- [33] Notion Labs, Inc. 2026. Notion AI Autofill for Databases. Notion Help Center, <https://www.notion.com/help/autofill>. Accessed 29 September 2026.
- [34] Azam Nouri. 2026. StepCache: Step-Level Reuse with Lightweight Verification and Selective Patching for LLM Serving. arXiv:2603.28795
- [35] Shuyin Ouyang, Jie M. Zhang, Mark Harman, and Meng Wang. 2025. An Empirical Study of the Non-Determinism of ChatGPT in Code Generation. *ACM Trans. Softw. Eng. Methodol.* 34, 2 (2025), 42:1–42:28. <https://doi.org/10.1145/3697010> arXiv:2308.02828.
- [36] Nicholas Pangakis, Samuel Wolken, and Neil Fasching. 2023. Automated Annotation with Generative AI Requires Validation. arXiv:2306.00176
- [37] Liana Patel, Siddharth Jha, Melissa Pan, Harshit Gupta, Parth Asawa, Carlos Guestrin, and Matei Zaharia. 2025. Semantic Operators and Their Optimization: Enabling LLM-Based Data Processing with Accuracy Guarantees in LOTUS. *Proc. VLDB Endow.* 18, 11 (2025), 4171–4184. <https://doi.org/10.14778/3749646.3749685>
- [38] Luis Gaspar Schroeder, Aditya Desai, Alejandro Cuadron, Kyle Chu, Shu Liu, Mark Zhao, Stephan Krusche, Alfons Kemper, Matei Zaharia, and Joseph E. Gonzalez. 2026. vCache: Verified Semantic Prompt Caching. In *Proc. Int. Conf. on Learning Representations (ICLR)*. arXiv:2502.03771.
- [39] R. J. Serfling. 1974. Probability Inequalities for the Sum in Sampling without Replacement. *Ann. Statist.* 2, 1 (1974), 39–48. <https://doi.org/10.1214/aos/1176342611>
- [40] Shreya Shankar, Tristan Chambers, Tarak Shah, Aditya G. Parameswaran, and Eugene Wu. 2025. DocETL: Agentic Query Rewriting and Evaluation for Complex Document Processing. *Proc. VLDB Endow.* 18, 9 (2025), 3035–3048. <https://doi.org/10.14778/3746405.3746426>
- [41] Shreya Shankar, Haotian Li, Parth Asawa, Madelon Hulsebos, Yiming Lin, J. D. Zambrescu-Pereira, Harrison Chase, Will Fu-Hinthorn, Aditya G. Parameswaran, and Eugene Wu. 2024. SPADe: Synthesizing Data Quality Assertions for Large Language Model Pipelines. arXiv:2401.03038
- [42] Shreya Shankar, Sepanta Zeighami, and Aditya Parameswaran. 2026. Task Cascades for Efficient Unstructured Data Processing. *Proc. ACM Manag. Data* (2026). <https://doi.org/10.1145/3786702> SIGMOD 2026.
- [43] Snowflake Inc. 2026. Snowflake Cortex AI Functions: AI\_COMPLETE and AI\_CLASSIFY. Snowflake Documentation, [https://docs.snowflake.com/en/sql-reference/functions/ai\\_complete](https://docs.snowflake.com/en/sql-reference/functions/ai_complete). Accessed 29 September 2026.
- [44] Petter Törnberg. 2023. ChatGPT-4 Outperforms Experts and Crowd Workers in Annotating Political Twitter Messages with Zero-Shot Learning. arXiv:2304.06588
- [45] Ruodu Wang and Aaditya Ramdas. 2022. False Discovery Rate Control with E-values. *J. R. Stat. Soc. Series B* 84, 3 (2022), 822–852. <https://doi.org/10.1111/rssb.12489>
- [46] Ian Waudby-Smith and Aaditya Ramdas. 2024. Estimating Means of Bounded Random Variables by Betting. *J. R. Stat. Soc. Series B* 86, 1 (2024), 1–27. <https://doi.org/10.1093/jrsssb/qqad009>
- [47] Jiayi Yuan, Hao Li, Xinheng Ding, Wenya Xie, Yu-Jhe Li, Wentian Zhao, Kun Wan, Jing Shi, Xia Hu, and Zirui Liu. 2025. Understanding and Mitigating Numerical Sources of Nondeterminism in LLM Inference. In *Advances in Neural Information Processing Systems (NeurIPS)*. arXiv:2506.09501.
- [48] Sepanta Zeighami, Shreya Shankar, and Aditya Parameswaran. 2026. Cut Costs, Not Accuracy: LLM-Powered Data Processing with Guarantees. *Proc. ACM Manag. Data* (2026). <https://doi.org/10.1145/3769776> SIGMOD 2026.
- [49] Caleb Ziem, William Held, Omar Shaikh, Jiaao Chen, Zhehao Zhang, and Diyi Yang. 2024. Can Large Language Models Transform Computational Social Science? *Comput. Linguist.* 50, 1 (2024), 237–291. [https://doi.org/10.1162/coli\\_a\\_00502](https://doi.org/10.1162/coli_a_00502)